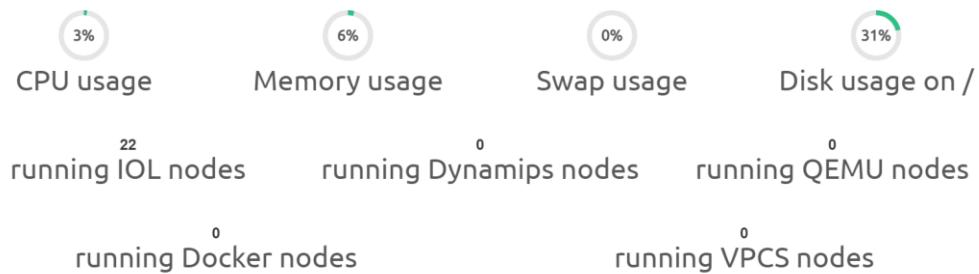
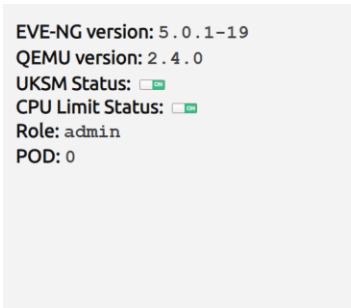
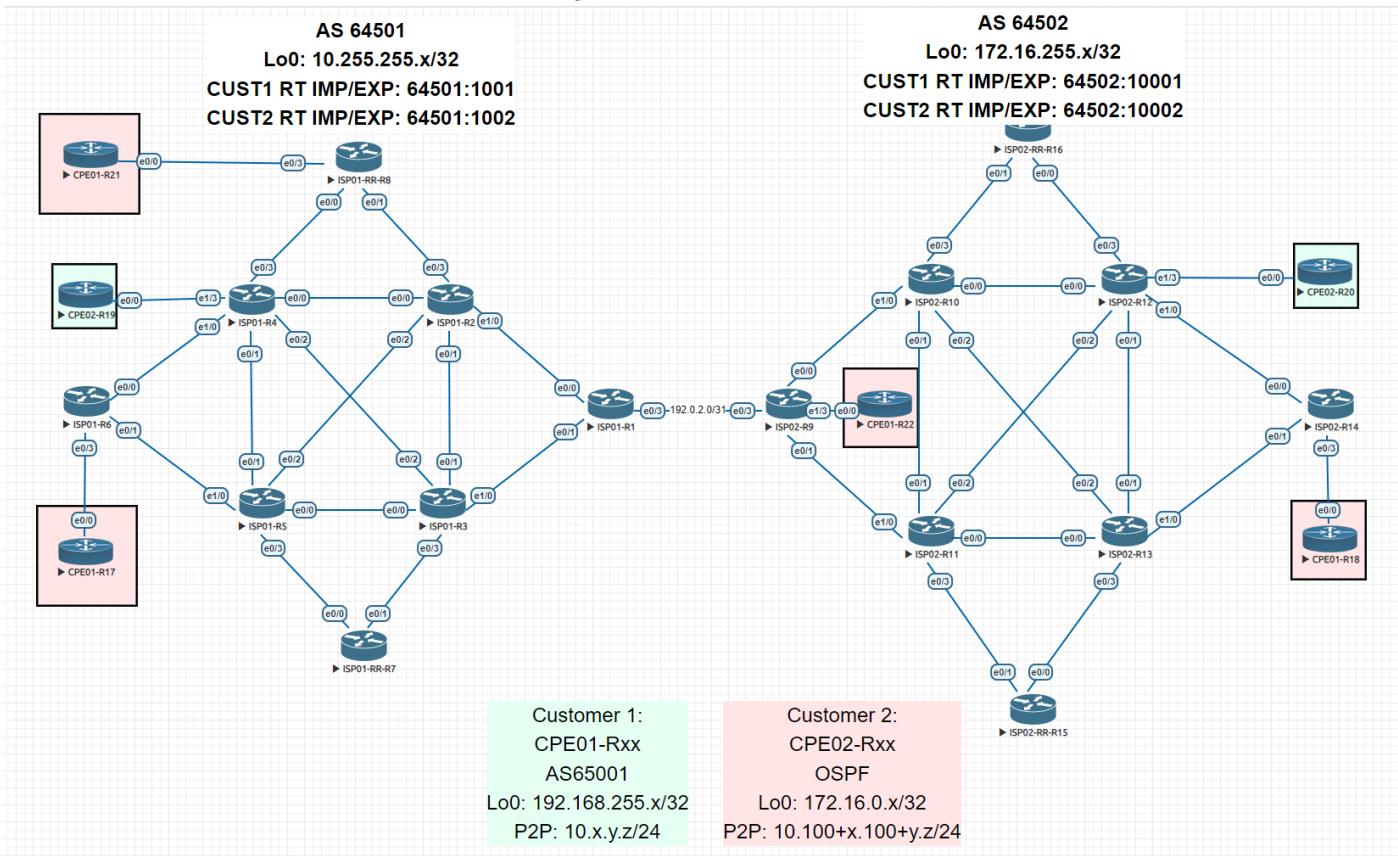


# Cisco eBGP option B – MPLS L3VPN



Использованные образы:

- i86bi\_LinuxL3-AdvEnterpriseK9-M2\_157\_3\_May\_2018.bin

## **0. Summary**

В данной лабораторной работе необходимо организовать связность между сайтами клиентов, подключенных в разных AS. Для этого необходимо организовать interAS-стык по схеме «option B», т.е. поднятие между ASBR'ми interAS eBGP-сессии в AFI/SAFI VPNV4 Unicast.

## 1. Solution

По умолчанию, на оборудовании Cisco фильтруются получаемые VPNv4-маршруты, если, конечно, получает не рефлектор. Т.к. ASBR'ы у нас не являются рефлекторами, отключаем фильтрацию при помощи команды no bgp default route-target filter.

```
ISP01-ASBR-R1#show bgp vpnv4 uni all | b Network
      Network          Next Hop          Metric LocPrf Weight Path
Route Distinguisher: 10.255.255.1:11 (default for vrf MGMT)
  *> 10.0.0.1/32    0.0.0.0          0       32768 ?
  *>i 10.0.0.2/32  10.255.255.2    0       100     0 ?
  *>i 10.0.0.3/32  10.255.255.3    0       100     0 ?
  *>i 10.0.0.4/32  10.255.255.4    0       100     0 ?
  *>i 10.0.0.5/32  10.255.255.5    0       100     0 ?
  *>i 10.0.0.6/32  10.255.255.6    0       100     0 ?
  *>i 10.0.0.7/32  10.255.255.7    0       100     0 ?
  *>i 10.0.0.8/32  10.255.255.8    0       100     0 ?
Route Distinguisher: 10.255.255.2:11
  * i 10.0.0.2/32  10.255.255.2    0       100     0 ?
  *>i               10.255.255.2    0       100     0 ?
Route Distinguisher: 10.255.255.3:11
  * i 10.0.0.3/32  10.255.255.3    0       100     0 ?
  *>i               10.255.255.3    0       100     0 ?
Route Distinguisher: 10.255.255.4:11
  * i 10.0.0.4/32  10.255.255.4    0       100     0 ?
  *>i               10.255.255.4    0       100     0 ?
Route Distinguisher: 10.255.255.5:11
  * i 10.0.0.5/32  10.255.255.5    0       100     0 ?
  *>i               10.255.255.5    0       100     0 ?
Route Distinguisher: 10.255.255.6:11
  Network          Next Hop          Metric LocPrf Weight Path
  * i 10.0.0.6/32  10.255.255.6    0       100     0 ?
  *>i               10.255.255.6    0       100     0 ?
Route Distinguisher: 10.255.255.7:11
  * i 10.0.0.7/32  10.255.255.7    0       100     0 ?
  *>i               10.255.255.7    0       100     0 ?
Route Distinguisher: 10.255.255.8:11
  * i 10.0.0.8/32  10.255.255.8    0       100     0 ?
  *>i               10.255.255.8    0       100     0 ?
ISP01-ASBR-R1#
```

После применения:

```
ISP01-ASBR-R1#show bgp vpnv4 uni all | b Network
      Network          Next Hop          Metric LocPrf Weight Path
Route Distinguisher: 10.255.255.1:11 (default for vrf MGMT)
  *> 10.0.0.1/32    0.0.0.0          0       32768 ?
  *>i 10.0.0.2/32  10.255.255.2    0       100     0 ?
  *>i 10.0.0.3/32  10.255.255.3    0       100     0 ?
  *>i 10.0.0.4/32  10.255.255.4    0       100     0 ?
  *>i 10.0.0.5/32  10.255.255.5    0       100     0 ?
  *>i 10.0.0.6/32  10.255.255.6    0       100     0 ?
  *>i 10.0.0.7/32  10.255.255.7    0       100     0 ?
  *>i 10.0.0.8/32  10.255.255.8    0       100     0 ?
Route Distinguisher: 10.255.255.2:11
  * i 10.0.0.2/32  10.255.255.2    0       100     0 ?
  *>i               10.255.255.2    0       100     0 ?
Route Distinguisher: 10.255.255.3:11
  * i 10.0.0.3/32  10.255.255.3    0       100     0 ?
  *>i               10.255.255.3    0       100     0 ?
Route Distinguisher: 10.255.255.4:11
  * i 10.0.0.4/32  10.255.255.4    0       100     0 ?
  *>i               10.255.255.4    0       100     0 ?
Route Distinguisher: 10.255.255.4:1002
  * i 10.104.119.0/24 10.255.255.4   0       100     0 ?
  *>i               10.255.255.4   0       100     0 ?
  * i 172.16.0.19/32 10.255.255.4  11      100     0 ?
  *>i               10.255.255.4  11      100     0 ?
Route Distinguisher: 10.255.255.5:11
  * i 10.0.0.5/32  10.255.255.5    0       100     0 ?
  *>i               10.255.255.5    0       100     0 ?
Route Distinguisher: 10.255.255.6:11
  * i 10.0.0.6/32  10.255.255.6    0       100     0 ?
  *>i               10.255.255.6    0       100     0 ?
Route Distinguisher: 10.255.255.6:1001
  * i 10.6.17.0/24  10.255.255.6    0       100     0 ?
  *>i               10.255.255.6    0       100     0 ?
  * i 192.168.255.17/32
```

```

          10.255.255.6      0   100   0 65001 i
 *>i      10.255.255.6      0   100   0 65001 i
Route Distinguisher: 10.255.255.7:11
 * i 10.0.0.7/32      10.255.255.7      0   100   0 ?
 *>i      10.255.255.7      0   100   0 ?
Route Distinguisher: 10.255.255.8:11
 * i 10.0.0.8/32      10.255.255.8      0   100   0 ?
 *>i      10.255.255.8      0   100   0 ?
Route Distinguisher: 10.255.255.8:1001
 *>i 10.8.21.0/24      10.255.255.8      0   100   0 ?
 * i      10.255.255.8      0   100   0 ?
 *>i 192.168.255.21/32      10.255.255.8      0   100   0 65001 i
 * i      10.255.255.8      0   100   0 65001 i

```

После этого настроим eBGP-стык между ASBR'ами.

| ISP01-ASBR-R1  | ISP02-ASBR-R09   |
|--|--|
| <pre> router bgp 64501  no bgp default route-target filter  neighbor 192.0.2.1 remote-as 64502  neighbor 192.0.2.1 update-source Ethernet0/3  !  address-family vpnv4  neighbor 192.0.2.1 activate  neighbor 192.0.2.1 send-community both  exit-address-family  ! !</pre> | <pre> router bgp 64502  no bgp default route-target filter  neighbor 192.0.2.0 remote-as 64501  neighbor 192.0.2.0 update-source Ethernet0/3  !  address-family vpnv4  neighbor 192.0.2.0 activate  neighbor 192.0.2.0 send-community both  exit-address-family  ! !</pre> |

После поднятия BGP-сессии на интерфейсах Ethernet0/3 автоматически будет включен MPLS:

```

ISP02-ASBR-R09#
*Mar 18 18:18:24.693: %BGP-5-ADJCHANGE: neighbor 192.0.2.0 Up
ISP02-ASBR-R09#
*Mar 18 18:18:24.698: %BGP_LMM-6-AUTOGN1: The mpls bgp forwarding command has been configured on
interface: Ethernet0/3
ISP02-ASBR-R09#show mpls interface
Interface          IP           Tunnel   BGP  Static Operational
Ethernet0/0        Yes (ldp)    No       No   No    Yes
Ethernet0/1        Yes (ldp)    No       No   No    Yes
Ethernet0/3        No          No       Yes  No    Yes
ISP02-ASBR-R09#

```

Проверим что сейчас в VPNv4 Unicast BGP Rib на ISP01-ASBR-R1:

```

ISP01-ASBR-R1#show bgp vpnv4 unicast all | b Network
 Network      Next Hop      Metric LocPrf Weight Path
Route Distinguisher: 10.255.255.1:11 (default for vrf MGMT)
 *> 10.0.0.1/32      0.0.0.0      0      32768 ?
 *>i 10.0.0.2/32      10.255.255.2      0      100   0 ?
 *>i 10.0.0.3/32      10.255.255.3      0      100   0 ?
 *>i 10.0.0.4/32      10.255.255.4      0      100   0 ?
 *>i 10.0.0.5/32      10.255.255.5      0      100   0 ?
 *>i 10.0.0.6/32      10.255.255.6      0      100   0 ?
 *>i 10.0.0.7/32      10.255.255.7      0      100   0 ?
 *>i 10.0.0.8/32      10.255.255.8      0      100   0 ?
Route Distinguisher: 10.255.255.2:11
 * i 10.0.0.2/32      10.255.255.2      0      100   0 ?
 *>i      10.255.255.2      0      100   0 ?
Route Distinguisher: 10.255.255.3:11
 * i 10.0.0.3/32      10.255.255.3      0      100   0 ?
 *>i      10.255.255.3      0      100   0 ?
Route Distinguisher: 10.255.255.4:11
 * i 10.0.0.4/32      10.255.255.4      0      100   0 ?
 *>i      10.255.255.4      0      100   0 ?
Route Distinguisher: 10.255.255.4:1002
 * i 10.104.119.0/24      10.255.255.4      0      100   0 ?
 *>i      10.255.255.4      0      100   0 ?
 * i 172.16.0.19/32      10.255.255.4      11     100   0 ?
 *>i      10.255.255.4      11     100   0 ?
Route Distinguisher: 10.255.255.5:11
 * i 10.0.0.5/32      10.255.255.5      0      100   0 ?
 *>i      10.255.255.5      0      100   0 ?
Route Distinguisher: 10.255.255.6:11

```

```

* i 10.0.0.6/32      10.255.255.6          0   100   0 ?
*+i                           10.255.255.6          0   100   0 ?
Route Distinguisher: 10.255.255.6:1001
* i 10.6.17.0/24     10.255.255.6          0   100   0 ?
*+i                           10.255.255.6          0   100   0 ?
* i 192.168.255.17/32
                           10.255.255.6          0   100   0 65001 i
*+i                           10.255.255.6          0   100   0 65001 i
Route Distinguisher: 10.255.255.7:11
* i 10.0.0.7/32     10.255.255.7          0   100   0 ?
*+i                           10.255.255.7          0   100   0 ?
Route Distinguisher: 10.255.255.8:11
* i 10.0.0.8/32     10.255.255.8          0   100   0 ?
*+i                           10.255.255.8          0   100   0 ?
Route Distinguisher: 10.255.255.8:1001
*+i 10.8.21.0/24     10.255.255.8          0   100   0 ?
* i                           10.255.255.8          0   100   0 ?
*+i 192.168.255.21/32
                           10.255.255.8          0   100   0 65001 i
* i                           10.255.255.8          0   100   0 65001 i
Route Distinguisher: 172.16.255.9:11
*+> 10.0.0.9/32     192.0.2.1          0       0 64502 ?
Route Distinguisher: 172.16.255.9:10001
*+> 10.9.22.0/24     192.0.2.1          0       0 64502 ?
*+> 192.168.255.22/32
                           192.0.2.1          0 64502 65001 i
Route Distinguisher: 172.16.255.10:11
*+> 10.0.0.10/32     192.0.2.1          0       0 64502 ?
Route Distinguisher: 172.16.255.11:11
*+> 10.0.0.11/32     192.0.2.1          0       0 64502 ?
Route Distinguisher: 172.16.255.12:11
*+> 10.0.0.12/32     192.0.2.1          0       0 64502 ?
Route Distinguisher: 172.16.255.12:10002
*+> 10.112.120.0/24   192.0.2.1          0       0 64502 ?
*+> 172.16.0.20/32   192.0.2.1          0       0 64502 ?
Route Distinguisher: 172.16.255.13:11
*+> 10.0.0.13/32     192.0.2.1          0       0 64502 ?
Route Distinguisher: 172.16.255.14:11
*+> 10.0.0.14/32     192.0.2.1          0       0 64502 ?
Route Distinguisher: 172.16.255.14:10001
*+> 10.14.18.0/24     192.0.2.1          0       0 64502 ?
*+> 192.168.255.18/32
                           192.0.2.1          0 64502 65001 i
Route Distinguisher: 172.16.255.15:11
*+> 10.0.0.15/32     192.0.2.1          0       0 64502 ?
Route Distinguisher: 172.16.255.16:11
*+> 10.0.0.16/32     192.0.2.1          0       0 64502 ?
ISP01-ASBR-R1#

```

Здесь мы видим, что от другого ASBR мы получили необходимую информацию.

Теперь проверим, что у нас есть на рефлекторах:

```

ISP01-RR-R8#show bgp vpng4 uni all | b Route Distinguisher: 172.16.255
Route Distinguisher: 172.16.255.9:11
* i 10.0.0.9/32     192.0.2.1          0   100   0 64502 ?
Route Distinguisher: 172.16.255.9:10001
* i 10.9.22.0/24   192.0.2.1          0   100   0 64502 ?
* i 192.168.255.22/32
                           192.0.2.1          0   100   0 64502 65001 i
Route Distinguisher: 172.16.255.10:11
* i 10.0.0.10/32   192.0.2.1          0   100   0 64502 ?
Route Distinguisher: 172.16.255.11:11
* i 10.0.0.11/32   192.0.2.1          0   100   0 64502 ?
Route Distinguisher: 172.16.255.12:11
* i 10.0.0.12/32   192.0.2.1          0   100   0 64502 ?
Route Distinguisher: 172.16.255.12:10002
* i 10.112.120.0/24 192.0.2.1          0   100   0 64502 ?
* i 172.16.0.20/32 192.0.2.1          0   100   0 64502 ?
Route Distinguisher: 172.16.255.13:11
* i 10.0.0.13/32   192.0.2.1          0   100   0 64502 ?
Route Distinguisher: 172.16.255.14:11
* i 10.0.0.14/32   192.0.2.1          0   100   0 64502 ?
Route Distinguisher: 172.16.255.14:10001
* i 10.14.18.0/24   192.0.2.1          0   100   0 64502 ?
* i 192.168.255.18/32
                           192.0.2.1          0   100   0 64502 65001 i
      Network          Next Hop          Metric LocPrf Weight Path
Route Distinguisher: 172.16.255.15:11
* i 10.0.0.15/32   192.0.2.1          0   100   0 64502 ?

```

```
Route Distinguisher: 172.16.255.16:11
```

```
* i 10.0.0.16/32      192.0.2.1          0   100      0 64502 ?
```

Мы видим, что на рефлекторы необходимая маршрутная информация попала. Теперь проверим наличие в RIB для vrf CUST1 наличие нужных маршрутов:

```
ISP01-RR-R8#show ip route vrf CUST1 | b Gateway  
Gateway of last resort is not set
```

```
10.0.0.0/8 is variably subnetted, 3 subnets, 2 masks  
B      10.6.17.0/24 [200/0] via 10.255.255.6, 00:27:35  
C      10.8.21.0/24 is directly connected, Ethernet0/3  
L      10.8.21.8/32 is directly connected, Ethernet0/3  
      192.168.255.0/32 is subnetted, 2 subnets  
B      192.168.255.17 [200/0] via 10.255.255.6, 00:27:35  
B      192.168.255.21 [20/0] via 10.8.21.21, 00:27:35  
ISP01-RR-R8#show bgp vpnv4 unicast vrf CUST1 | b Network  
      Network          Next Hop          Metric LocPrf Weight Path  
Route Distinguisher: 10.255.255.8:1001 (default for vrf CUST1)  
*>i 10.6.17.0/24      10.255.255.6      0     100      0 ?  
* > 10.8.21.0/24      0.0.0.0          0           32768 ?  
*>i 192.168.255.17/32  
      10.255.255.6      0     100      0 65001 i  
*> 192.168.255.21/32  
      10.8.21.21        0           0 65001 i  
ISP01-RR-R8#
```

И здесь мы видим, что в RIB vrf CUST1 ничего не попало, что, впрочем, ожидаемо, т.к. route-target у ISP01 и ISP02 для этого vrf отличаются. На маршрутизаторах ISP01, к которым подключены СРЕ01, добавим RT на импорт и экспорт 64502:10001:

| ISP01-RR-R8  | ISP01-R6   |
|--|--|
| vrf definition CUST1<br>!<br>address-family ipv4<br>route-target export 64502:10001<br>route-target import 64502:10001<br>exit-address-family<br>! | vrf definition CUST1<br>!<br>address-family ipv4<br>route-target export 64502:10001<br>route-target import 64502:10001<br>exit-address-family<br>! |

Проверим еще раз, и увидим, что в BGP RIB на RR маршруты попали, а вот в RIB уже нет:

```
ISP01-RR-R8#show ip route vrf CUST1 | b Gateway  
Gateway of last resort is not set
```

```
10.0.0.0/8 is variably subnetted, 3 subnets, 2 masks  
B      10.6.17.0/24 [200/0] via 10.255.255.6, 00:33:50  
C      10.8.21.0/24 is directly connected, Ethernet0/3  
L      10.8.21.8/32 is directly connected, Ethernet0/3  
      192.168.255.0/32 is subnetted, 2 subnets  
B      192.168.255.17 [200/0] via 10.255.255.6, 00:33:50  
B      192.168.255.21 [20/0] via 10.8.21.21, 00:00:30  
ISP01-RR-R8#show bgp vpnv4 unicast vrf CUST1 | b Network  
      Network          Next Hop          Metric LocPrf Weight Path  
Route Distinguisher: 10.255.255.8:1001 (default for vrf CUST1)  
*>i 10.6.17.0/24      10.255.255.6      0     100      0 ?  
* > 10.8.21.0/24      0.0.0.0          0           32768 ?  
* i 10.9.22.0/24      192.0.2.1        0     100      0 64502 ?  
* i 10.14.18.0/24      192.0.2.1        0     100      0 64502 ?  
*>i 192.168.255.17/32  
      10.255.255.6      0     100      0 65001 i  
* i 192.168.255.18/32  
      192.0.2.1          0     100      0 64502 65001 i  
*> 192.168.255.21/32  
      10.8.21.21        0           0 65001 i  
* i 192.168.255.22/32  
      192.0.2.1          0     100      0 64502 65001 i  
ISP01-RR-R8#
```

Посмотрим более внимательно, что находится в BGP RIB на одном из полученных маршрутов (спойлер, ответ там будет):

```

ISP01-RR-R8#show bgp vpng4 unicast vrf CUST1 192.168.255.18/32
BGP routing table entry for 10.255.255.8:1001:192.168.255.18/32, version 0
Paths: (1 available, no best path)
    Not advertised to any peer
    Refresh Epoch 2
64502 65001, (Received from a RR-client), imported safety path from 172.16.255.14:10001:192.168.255.18/32 (global)
    192.0.2.1 (inaccessible) (via default) from 10.255.255.1 (10.255.255.1)
        Origin IGP, metric 0, localpref 100, valid, internal
        Extended Community: RT:64502:10001
        mpls labels in/out nolabel/90018
        rx pathid: 0, tx pathid: 0
ISP01-RR-R8#

```

Здесь мы видим, что маршрут признан некорректным из-за того, что нет маршрута до next-hop'a 192.0.2.1. Собственно, есть два варианта решения данной проблемы:

1. Использовать next-hop-self на ASBR'e в сторону рефлекторов.
2. Добавить адрес next-hop в IGP.

Наиболее распространенным и простым является использование next-hop-self. Именно им мы воспользуемся в AS64501:

|   |
|---|
| ISP01-ASBR-R1   |
| <pre> router bgp 64501 ! address-family vpng4     neighbor RR send-community both     neighbor RR next-hop-self exit-address-family !</pre> |

Проверим снова на RR:

```

ISP01-RR-R8#show bgp vpng4 unicast vrf CUST1 192.168.255.18/32
BGP routing table entry for 10.255.255.8:1001:192.168.255.18/32, version 70
Paths: (1 available, best #1, table CUST1)
    Advertised to update-groups:
        1
    Refresh Epoch 5
64502 65001, (Received from a RR-client), imported path from 172.16.255.14:10001:192.168.255.18/32 (global)
    10.255.255.1 (metric 200) (via default) from 10.255.255.1 (10.255.255.1)
        Origin IGP, metric 0, localpref 100, valid, internal, best
        Extended Community: RT:64502:10001
        mpls labels in/out nolabel/10031
        rx pathid: 0, tx pathid: 0x0
ISP01-RR-R8#

```

И на ISP01-R6:

```

ISP01-R6#show bgp vpng4 unicast vrf CUST1 | b Network
      Network          Next Hop          Metric LocPrf Weight Path
Route Distinguisher: 10.255.255.6:1001 (default for vrf CUST1)
* > 10.6.17.0/24    0.0.0.0            0          32768 ?
* >i 10.8.21.0/24   10.255.255.8       0          100     0 ?
* >i 10.9.22.0/24   10.255.255.1       0          100     0 64502 ?
* >i 10.14.18.0/24  10.255.255.1       0          100     0 64502 ?
* > 192.168.255.17/32
                  10.6.17.17          0          0 65001 i
* >i 192.168.255.18/32
                  10.255.255.1       0          100     0 64502 65001 i
* >i 192.168.255.21/32
                  10.255.255.8       0          100     0 65001 i
* >i 192.168.255.22/32
                  10.255.255.1       0          100     0 64502 65001 i
ISP01-R6#show ip route vrf CUST1 | b Gateway
Gateway of last resort is not set

      10.0.0.0/8 is variably subnetted, 5 subnets, 2 masks
C        10.6.17.0/24 is directly connected, Ethernet0/3
L        10.6.17.6/32 is directly connected, Ethernet0/3
B        10.8.21.0/24 [200/0] via 10.255.255.8, 00:18:25
B        10.9.22.0/24 [200/0] via 10.255.255.1, 00:01:08
B        10.14.18.0/24 [200/0] via 10.255.255.1, 00:01:08
      192.168.255.0/32 is subnetted, 4 subnets
B          192.168.255.17 [20/0] via 10.6.17.17, 00:01:08

```

```
B      192.168.255.18 [200/0] via 10.255.255.1, 00:01:08
B      192.168.255.21 [200/0] via 10.255.255.8, 00:18:53
B      192.168.255.22 [200/0] via 10.255.255.1, 00:01:08
ISP01-R6#
```

В общем-то, на стороне ISP01 мы закончили, однако, на данный момент связности между сайтами клиента в разных ISP нет. Поэтому идем настраивать вторую половину.

В рамках настройки ISP02 мы попробуем уже второй вариант решения проблемы с недоступным next-hop, а именно, добавим его в IGP (в данном случае, OSPFv2).

Прежде чем начать, посмотрим внимательно RIB на ISP02-ASBR-R09:

```
ISP02-ASBR-R09#show ip route | i 192.0.2.
  192.0.2.0/24 is variably subnetted, 5 subnets, 2 masks
C      192.0.2.0/31 is directly connected, Ethernet0/3
C      192.0.2.0/32 is directly connected, Ethernet0/3
L      192.0.2.1/32 is directly connected, Ethernet0/3
C      192.0.2.2/31 is directly connected, Ethernet0/3.10
L      192.0.2.3/32 is directly connected, Ethernet0/3.10
ISP02-ASBR-R09#
```

Ethernet0/3.10 – это временный сабинтерфейс, единственная цель которого показать разницу.

В выводе выше мы видим следующее: в RIB для Ethernet0/3 находится 3 записи:

- прописанный на интерфейсе local-префикс 192.0.2.1/32,
- connected-префикс 192.0.2.0/31,
- connected-префикс 192.0.2.0/32

А вот для E0/3.10 такого нет: у него только connected /31 и local /32.

Давайте сравним конфиг этих интерфейсов:

```
!
interface Ethernet0/3
  ip address 192.0.2.1 255.255.255.254
  duplex auto
  mpls bgp forwarding
!
interface Ethernet0/3.10
  encapsulation dot1Q 10
  ip address 192.0.2.3 255.255.255.254
!
```

mpls bgp forwarding – это команда, которая была добавлена автоматически при поднятии VPNv4-сессии между ASBR. В этот же момент, IOS/IOS-XE автоматически добавил /32 connected-маршрут: он необходим для корректной работы LSP, требующей резолва через /32-префикс. Кстати, в IOS-XR такого действия по умолчанию нет, поэтому приходится костылить со статиком /32.

Возвращаемся к добавлению адреса next-hop в IGP. У нас есть несколько вариантов:

1. Включение OSPF на интерфейсе в сторону другого ASBR (не забываем включить passive-interface).
2. Редистрибуция в OSPF (в нашем случае connected).

А теперь вспоминаем требование о необходимости резолва через /32: исходя из этого, единственный рабочий вариант для нас – это редистрибуция connected-маршрутов, причем с фильтрацией (см. вывод ниже).

```
ISP02-ASBR-R09(config)#router ospf 1
ISP02-ASBR-R09(config-router)#redistribute connected subnets
ISP02-ASBR-R09#
*Mar 19 06:20:32.647: %SYS-5-CONFIG_I: Configured from console by console
ISP02-ASBR-R09#
*Mar 19 06:20:32.652: %OSPF-4-CONFLICTING_LSAID: Process 1 area dummy area: LSA origination prevented by LSA with same
LSID but a different mask
  Existing Type 5 LSA: LSID 192.0.2.0/31
  New Destination: 192.0.2.0/32
ISP02-ASBR-R09#show ip ospf database external

  OSPF Router with ID (172.16.255.9) (Process ID 1)

    Type-5 AS External Link States
```

```

LS age: 77
Options: (No TOS-capability, DC, Upward)
LS Type: AS External Link
Link State ID: 192.0.2.0 (External Network Number )
Advertising Router: 172.16.255.9
LS Seq Number: 80000001
Checksum: 0x39DF
Length: 36
Network Mask: /31
    Metric Type: 2 (Larger than any link state path)
    MTID: 0
    Metric: 20
    Forward Address: 0.0.0.0
    External Route Tag: 0

```

ISP02-ASBR-R09#

Собственно, из вывода выше видно, что при обычной редистрибуции, в OSPF DB попадает только /31 префикс. Поэтому настраиваем фильтрацию:

```

ISP02-ASBR-R09
ip prefix-list PL_RDSTB_CON_TO OSPF seq 10 permit 192.0.2.0/32
!
route-map RM_RDSTB_CON_TO OSPF permit 10
  match ip address prefix-list PL_RDSTB_CON_TO OSPF
!
router ospf 1
  redistribute connected subnets route-map RM_RDSTB_CON_TO OSPF
!
```

ISP02-ASBR-R09#show ip ospf database external

OSPF Router with ID (172.16.255.9) (Process ID 1)

Type-5 AS External Link States

```

LS age: 6
Options: (No TOS-capability, DC, Upward)
LS Type: AS External Link
Link State ID: 192.0.2.0 (External Network Number )
Advertising Router: 172.16.255.9
LS Seq Number: 80000003
Checksum: 0x3BDA
Length: 36
Network Mask: /32
    Metric Type: 2 (Larger than any link state path)
    MTID: 0
    Metric: 20
    Forward Address: 0.0.0.0
    External Route Tag: 0

```

ISP02-ASBR-R09#

Проверим на ISP02-R14 и на CPE01-R18:

ISP02-R14#show ip route vrf CUST1 | b Gateway  
Gateway of last resort is not set

```

10.0.0.0/8 is variably subnetted, 5 subnets, 2 masks
B      10.6.17.0/24 [200/0] via 192.0.2.0, 00:00:20
B      10.8.21.0/24 [200/0] via 192.0.2.0, 00:07:22
B      10.9.22.0/24 [200/0] via 172.16.255.9, 00:39:16
C      10.14.18.0/24 is directly connected, Ethernet0/3
L      10.14.18.14/32 is directly connected, Ethernet0/3
      192.168.255.0/32 is subnetted, 4 subnets
B      192.168.255.17 [200/0] via 192.0.2.0, 00:00:20
B      192.168.255.18 [200/0] via 10.14.18.18, 00:39:28
B      192.168.255.21 [200/0] via 192.0.2.0, 00:07:22
B      192.168.255.22 [200/0] via 172.16.255.9, 00:39:16

```

ISP02-R14#show bgp vpnv4 unicast vrf CUST1 | b Network  
Network Next Hop Metric LocPrf Weight Path  
Route Distinguisher: 172.16.255.14:10001 (default for vrf CUST1)  
\*>i 10.6.17.0/24 192.0.2.0 0 100 0 64501 ?  
\*>i 10.8.21.0/24 192.0.2.0 0 100 0 64501 ?  
\*>i 10.9.22.0/24 172.16.255.9 0 100 0 ?

```

*> 10.14.18.0/24 0.0.0.0          0      32768 ?
*>i 192.168.255.17/32
    192.0.2.0          0      100      0 64501 65001 i
*> 192.168.255.18/32
    10.14.18.18          0      0 65001 i
*>i 192.168.255.21/32
    192.0.2.0          0      100      0 64501 65001 i
*>i 192.168.255.22/32
    172.16.255.9        0      100      0 65001 i
ISP02-R14#

```

```

CPE01-R18#show ip route | b Gateway
Gateway of last resort is not set

 10.0.0.0/8 is variably subnetted, 5 subnets, 2 masks
B     10.6.17.0/24 [20/0] via 10.14.18.14, 00:03:30
B     10.8.21.0/24 [20/0] via 10.14.18.14, 00:10:31
B     10.9.22.0/24 [20/0] via 10.14.18.14, 00:42:25
C     10.14.18.0/24 is directly connected, Ethernet0/0
L     10.14.18.18/32 is directly connected, Ethernet0/0
    192.168.255.0/32 is subnetted, 4 subnets
B       192.168.255.17 [20/0] via 10.14.18.14, 00:03:30
C       192.168.255.18 is directly connected, Loopback0
B       192.168.255.21 [20/0] via 10.14.18.14, 00:10:31
B       192.168.255.22 [20/0] via 10.14.18.14, 00:42:25
CPE01-R18#

```

Судя по выводу, control-plane у нас сошелся. Проверим наличие связности:

```

CPE01-R18#ping 192.168.255.21 so Lo0
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 192.168.255.21, timeout is 2 seconds:
Packet sent with a source address of 192.168.255.18
.....
Success rate is 0 percent (0/5)
CPE01-R18#traceroute 192.168.255.21 so Lo0
Type escape sequence to abort.
Tracing the route to 192.168.255.21
VRF info: (vrf in name/id, vrf out name/id)
 1 10.14.18.14 [AS 64502] 1 msec 1 msec 0 msec
 2 *
CPE01-R18#

```

Давайте разбираться, почему не проходит. Для начала посмотрим, есть ли MPLS-метка для префикса:

```

ISP02-R14#show bgp vpnv4 unicast vrf CUST1 192.168.255.21/32
BGP routing table entry for 172.16.255.14:10001:192.168.255.21/32, version 40
Paths: (1 available, best #1, table CUST1)
Flag: 0x100
    Advertised to update-groups:
        1
    Refresh Epoch 1
64501 65001, imported path from 10.255.255.8:1001:192.168.255.21/32 (global)
    192.0.2.0 (metric 20) (via default) from 172.16.255.15 (172.16.255.15)
        Origin IGP, metric 0, localpref 100, valid, internal, best
        Extended Community: RT:64501:1001 RT:64502:10001
        Originator: 172.16.255.9, Cluster list: 172.16.255.255
        mpls labels in/out nolabel/10028
        rx pathid: 0, tx pathid: 0x0
ISP02-R14#

```

Метка (10028) есть. На всякий случай убедимся в том, что в RIB есть маршрут до Next-hop'a.

```

ISP02-R14#show ip route 192.0.2.0
Routing entry for 192.0.2.0/32, 1 known subnets
O E2    192.0.2.0 [110/20] via 172.31.34.13, 00:19:31, Ethernet0/1
                                [110/20] via 172.31.24.12, 00:19:31, Ethernet0/0
ISP02-R14#

```

Проверим LFIB на наличие метки для адреса next-hop'a:

```

ISP02-R14#show mpls forwarding-table 192.0.2.0 32
Local      Outgoing      Prefix          Bytes Label      Outgoing      Next Hop
Label      Label         or Tunnel Id   Switched      interface
140001    No Label     192.0.2.0/32    0            Et0/0        172.31.24.12
                                         No Label     192.0.2.0/32    0            Et0/1        172.31.34.13

```

Собственно, мы видим, что метки нет, что и является проблемой. Также это можно было узнать и посмотрев FIB для vrf CUST1:

```
ISP02-R14#show ip cef vrf CUST1 192.168.255.21/32 detail
192.168.255.21/32, epoch 0, flags [rib defined all labels]
  recursive via 192.0.2.0 label 10028
    nexthop 172.31.24.12 Ethernet0/0 unusable: no label
    nexthop 172.31.34.13 Ethernet0/1 unusable: no label
ISP02-R14#
```

Метка не была выделена из-за настроек MPLS LDP на всех маршрутизаторах ISP02:

```
ISP02-R14#show run | i mpls ldp advertise
no mpls ldp advertise-labels
mpls ldp advertise-labels for ACL_ISP02_LOOPBACKS
ISP02-R14#show ip access-lists ACL_ISP02_LOOPBACKS
Standard IP access list ACL_ISP02_LOOPBACKS
  10 permit 172.16.255.0, wildcard bits 0.0.0.255 (24 matches)
ISP02-R14#
```

Чтобы исправить, нужно добавить нужную строчку в ACL на каждом роутере ISP02:

|  |
|--|
| ISP02-ASBR-R09 / ISP02-R10 / ISP02-R11 / ISP02-R12 / ISP02-R13 / ISP02-R14 / ISP02-RR-R15 / ISP02-RR-R16 |
| ip access-list standard ACL_ISP02_LOOPBACKS<br>permit 192.0.2.0 0.0.0.0<br>!                             |

Добавив, проверяем LFIB на ISP02-R14:

```
ISP02-R14#show mpls forwarding-table 192.0.2.0 32
Local      Outgoing   Prefix          Bytes Label  Outgoing     Next Hop
Label      Label       or Tunnel Id  Switched   interface
140001    120001     192.0.2.0/32   0          Et0/0       172.31.24.12
           130001     192.0.2.0/32   0          Et0/1       172.31.34.13
```

Судя по выводу, сейчас все сошлось. Проверим работу на CPE01-R18:

```
CPE01-R18#ping 192.168.255.21 so Lo0
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 192.168.255.21, timeout is 2 seconds:
Packet sent with a source address of 192.168.255.18
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 3/4/6 ms
```

Теперь посмотрим, что у нас по трейсу:

```
CPE01-R18#traceroute 192.168.255.21 so Lo0
Type escape sequence to abort.
Tracing the route to 192.168.255.21
VRF info: (vrf in name/id, vrf out name/id)
  1 10.14.18.14 [AS 64502] 1 msec 1 msec 1 msec
  2 172.31.24.12 [MPLS: Labels 120001/10028 Exp 0] 4 msec 3 msec 4 msec
  3 172.31.112.10 [MPLS: Labels 100001/10028 Exp 0] 3 msec 4 msec 5 msec
  4 172.31.109.9 [MPLS: Labels 90019/10028 Exp 0] 4 msec 3 msec 4 msec
  5 192.0.2.0 [MPLS: Label 10028 Exp 0] 5 msec 4 msec 3 msec
  6 10.1.2.2 [MPLS: Labels 20008/80011 Exp 0] 4 msec 5 msec 3 msec
  7 10.8.21.8 [AS 64501] [MPLS: Labels 0/80011 Exp 0] 2 msec 2 msec 3 msec
  8 10.8.21.21 [AS 64501] 5 msec * 12 msec
CPE01-R18#
```

Краткое описание:

| № | Запись  | Комментарий   |
|---|---|---|
| 1 | 10.14.18.14 [AS 64502]                          | Пакет от CPE01-R18 до ISP02-R14   |
| 2 | 172.31.24.12 [MPLS: Labels 120001/10028 Exp 0]  | Пакет от ISP02-R14 до ISP02-R12:<br>PUSH на R14 10028 – VPN-метка для префикса 192.168.255.21/32<br>PUSH на R12 120001 – транспортная метка до 192.0.2.0/32 |
| 3 | 172.31.112.10 [MPLS: Labels 100001/10028 Exp 0] | Пакет от ISP02-R12 до ISP02-R10<br>SWAP на R12 метки 120001 на 100001   |
| 4 | 172.31.109.9 [MPLS: Labels 90019/10028 Exp 0]   | Пакет от ISP02-R10 до ISP02-ASBR-R09<br>SWAP на R10 метки 100001 на 90019   |
| 5 | 192.0.2.0 [MPLS: Label 10028 Exp 0]             | Пакет от ISP02-ASBR-R09 до ISP01-ASBR-R1 (см. описание)   |

|   |   |   |
|---|---|---|
| 6 | 10.1.2.2 [MPLS: Labels 20008/80011 Exp 0]         | Пакет от ISP01-ASBR-R1 до ISP01-R2:<br>SWAP на R1 10028 на 80001<br>PUSH на R1 20008 – транспортная метка до 10.255.255.8 |
| 7 | 10.8.21.8 [AS 64501] [MPLS: Labels 0/80011 Exp 0] | Пакет от ISP01-R2 до ISP01-RR-R8:<br>SWAP на R2 20008 на 0 (exp-null)   |
| 8 | 10.8.21.21 [AS 64501]                             | Пакет от ISP01-RR-R8 до CPE01-R21:<br>POP на R8 0<br>POP на R8 80011  |

Теперь посмотрим более подробно. Начнем с R14:

```
ISP02-R14#show bgp vpng4 unicast vrf CUST1 192.168.255.21/32
BGP routing table entry for 172.16.255.14:10001:192.168.255.21/32, version 40
Paths: (1 available, best #1, table CUST1)
Flag: 0x100
    Advertised to update-groups:
        1
    Refresh Epoch 1
64501 65001, imported path from 10.255.255.8:1001:192.168.255.21/32 (global)
    192.0.2.0 (metric 20) (via default) from 172.16.255.15 (172.16.255.15)
        Origin IGP, metric 0, localpref 100, valid, internal, best
        Extended Community: RT:64501:1001 RT:64502:10001
        Originator: 172.16.255.9, Cluster list: 172.16.255.255
        mpls labels in/out nolabel/10028
        rx pathid: 0, tx pathid: 0x0
ISP02-R14#
ISP02-R14#show mpls forwarding-table 192.0.2.0 32
Local      Outgoing   Prefix          Bytes Label  Outgoing     Next Hop
Label      Label       or Tunnel Id   Switched    interface
140001    120001     192.0.2.0/32   0           Et0/0       172.31.24.12
            130001     192.0.2.0/32   0           Et0/1       172.31.34.13
ISP02-R14#
ISP02-R14#show mpls forwarding-table vrf CUST1 192.168.255.21 32 detail
Local      Outgoing   Prefix          Bytes Label  Outgoing     Next Hop
Label      Label       or Tunnel Id   Switched    interface
None       10028      192.168.255.21/32[V]  \
                           Et0/0       172.31.24.12
    MAC/Encaps=14/22, MRU=1496, Label Stack{120001 10028}
    AABBCC00C001AABBCC00E0008847 1D4C10000272C000
    VPN route: CUST1
    No output feature configured
    Per-destination load-sharing, slots: 0
        10028      192.168.255.21/32[V]  \
                           Et0/1       172.31.34.13
    MAC/Encaps=14/22, MRU=1496, Label Stack{130001 10028}
    AABBCC00D001AABBCC00E0108847 1FB0D10000272C000
    VPN route: CUST1
    No output feature configured
    Per-destination load-sharing, slots: 1
ISP02-R14#
```

10028 – это VPN-метка, сгенерированная R1.

120001 и 130001 – это выходные метки до 192.0.2.0/32.

Соответственно, до префикса 192.168.255.21/32 возможны два стека меток, в зависимости от выбранного выходного интерфейса:

1. {120001 10028} – через R12
2. {130001 10028} – через R13

На транзитных маршрутизаторах ничего особо интересного не происходит, только swap меток. А вот на ASBR'ы нужно посмотреть внимательно:

Для начала, посмотрим, что именно находится за меткой 90019 в LFIB на R9:

```
ISP02-ASBR-R09#show mpls forwarding-table labels 90019
Local      Outgoing   Prefix          Bytes Label  Outgoing     Next Hop
Label      Label       or Tunnel Id   Switched    interface
90019     Pop Label  192.0.2.0/32   3088       Et0/3       192.0.2.0
ISP02-ASBR-R09#
```

Здесь мы явно видим выходной интерфейс (Ethernet0/3) и действие – снятие верхней метки. В трейсе, в пункте 5, мы это и видим: метка 90019 была снята и MPLS-пакет с меткой 10028 была отправлена через интерфейс Ethernet0/3 в сторону ISP01-ASBR-R1.

Теперь посмотрим, что внутри R1. Начнем с того, что скрывается за меткой 10028.

```
ISP01-ASBR-R1#show mpls forwarding-table labels 10028
Local      Outgoing   Prefix          Bytes Label  Outgoing   Next Hop
Label      Label       or Tunnel Id  Switched    interface
10028     80011      10.255.255.8:1001:192.168.255.21/32 \
                           2968           Et0/0        10.1.2.2
```

Судя по выводу, это VPN-метка для префикса 192.168.255.21/32 с RD 10.255.255.8:1001. Посмотрим на этот префикс в BGP RIB:

```
ISP01-ASBR-R1#show bgp vpnv4 unicast all 192.168.255.21/32
BGP routing table entry for 10.255.255.8:1001:192.168.255.21/32, version 38
Paths: (2 available, best #2, no table)
  Advertised to update-groups:
    2
  Refresh Epoch 2
65001
  10.255.255.8 (metric 200) (via default) from 10.255.255.7 (10.255.255.7)
    Origin IGP, metric 0, localpref 100, valid, internal
    Extended Community: RT:64501:1001 RT:64502:10001
    Originator: 10.255.255.8, Cluster list: 10.255.255.7
    mpls labels in/out 10028/80011
      rx pathid: 0, tx pathid: 0
  Refresh Epoch 2
65001
  10.255.255.8 (metric 200) (via default) from 10.255.255.8 (10.255.255.8)
    Origin IGP, metric 0, localpref 100, valid, internal, best
    Extended Community: RT:64501:1001 RT:64502:10001
    mpls labels in/out 10028/80011
      rx pathid: 0, tx pathid: 0x0
ISP01-ASBR-R1#
```

А метку для 10.255.255.8 (next-hop для указанного префикса) узнаем также через LFIB:

```
ISP01-ASBR-R1#show mpls forwarding-table 10.255.255.8 32
Local      Outgoing   Prefix          Bytes Label  Outgoing   Next Hop
Label      Label       or Tunnel Id  Switched    interface
10005     20008      10.255.255.8/32 0           Et0/0        10.1.2.2
ISP01-ASBR-R1#
```

Соответственно, label stack в сторону R8 будет {20008 80011}.

Также эту информацию можно посмотреть и так:

```
ISP01-ASBR-R1#show mpls forwarding-table labels 10028 detail
Local      Outgoing   Prefix          Bytes Label  Outgoing   Next Hop
Label      Label       or Tunnel Id  Switched    interface
10028     80011      10.255.255.8:1001:192.168.255.21/32 \
                           2968           Et0/0        10.1.2.2
  MAC/Encaps=14/22, MRU=1496, Label Stack{20008 80011}
  AABCC002001AABCC0010008847 04E280001388B000
  No output feature configured
ISP01-ASBR-R1#
```

В общем и целом, мы получили связность между сайтами клиента, подключенных к разным AS.

Ну и напоследок, поправим настройки для vrf CUST2 – теперь уже с обеих сторон:

| ISP01-R4   | ISP02-R12   |
|--|---|
| vrf definition CUST2 ! address-family ipv4   route-target export 64502:10002 exit-address-family ! | vrf definition CUST2 ! address-family ipv4   route-target export 64501:1002 exit-address-family ! |

### 3. Option B solution config:

AS 64501:

```
#ISP01-ASBR-R1:  
router bgp 64501  
no bgp default route-target filter  
neighbor 192.0.2.1 remote-as 64502  
neighbor 192.0.2.1 update-source Ethernet0/3  
!  
address-family vpnv4  
neighbor RR next-hop-self  
neighbor 192.0.2.1 activate  
neighbor 192.0.2.1 send-community both  
exit-address-family  
!  
!
```

```
#ISP01-RR-R8:  
!  
vrf definition CUST1  
!  
address-family ipv4  
route-target export 64502:10001  
route-target import 64502:10001  
exit-address-family  
!
```

```
#ISP01-R4:  
vrf definition CUST2  
!  
address-family ipv4  
route-target export 64502:10002  
exit-address-family  
!
```

```
#ISP01-R6:  
!  
vrf definition CUST1  
!  
address-family ipv4  
route-target export 64502:10001  
route-target import 64502:10001  
exit-address-family  
!
```

---

```
AS 64502:  
#ISP02-ASBR-R9:  
ip access-list standard ACL_ISP02_LOOPBACKS  
permit 192.0.2.0 0.0.0.0  
!  
ip prefix-list PL_RDSTB_CON_TO OSPF seq 10 permit 192.0.2.0/32  
!  
route-map RM_RDSTB_CON_TO OSPF permit 10  
match ip address prefix-list PL_RDSTB_CON_TO OSPF  
!  
router ospf 1  
redistribute connected subnets route-map RM_RDSTB_CON_TO OSPF  
!  
router bgp 64502  
no bgp default route-target filter  
neighbor 192.0.2.0 remote-as 64501  
neighbor 192.0.2.0 update-source Ethernet0/3  
!  
address-family vpnv4  
neighbor 192.0.2.0 activate  
neighbor 192.0.2.0 send-community both  
exit-address-family  
!
```

```
#ISP02-R10:  
ip access-list standard ACL_ISP02_LOOPBACKS  
permit 192.0.2.0 0.0.0.0  
!
```

```
#ISP02-R11:  
ip access-list standard ACL_ISP02_LOOPBACKS  
permit 192.0.2.0 0.0.0.0  
!
```

```
#ISP02-R12:  
ip access-list standard ACL_ISP02_LOOPBACKS  
permit 192.0.2.0 0.0.0.0  
!
```

```
vrf definition CUST2
!
address-family ipv4
  route-target export 64501:1002
exit-address-family
!

#ISP02-R13:
ip access-list standard ACL_ISP02_LOOPBACKS
  permit 192.0.2.0 0.0.0.0
!

#ISP02-R14:
ip access-list standard ACL_ISP02_LOOPBACKS
  permit 192.0.2.0 0.0.0.0
!

#ISP02-RR-R15:
ip access-list standard ACL_ISP02_LOOPBACKS
  permit 192.0.2.0 0.0.0.0
!

#ISP02-RR-R16:
ip access-list standard ACL_ISP02_LOOPBACKS
  permit 192.0.2.0 0.0.0.0
!
```