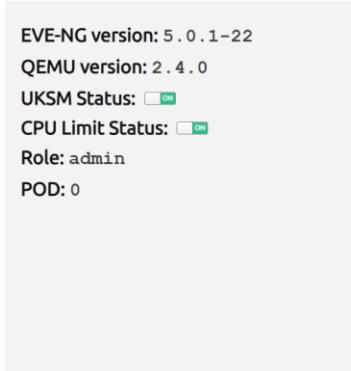
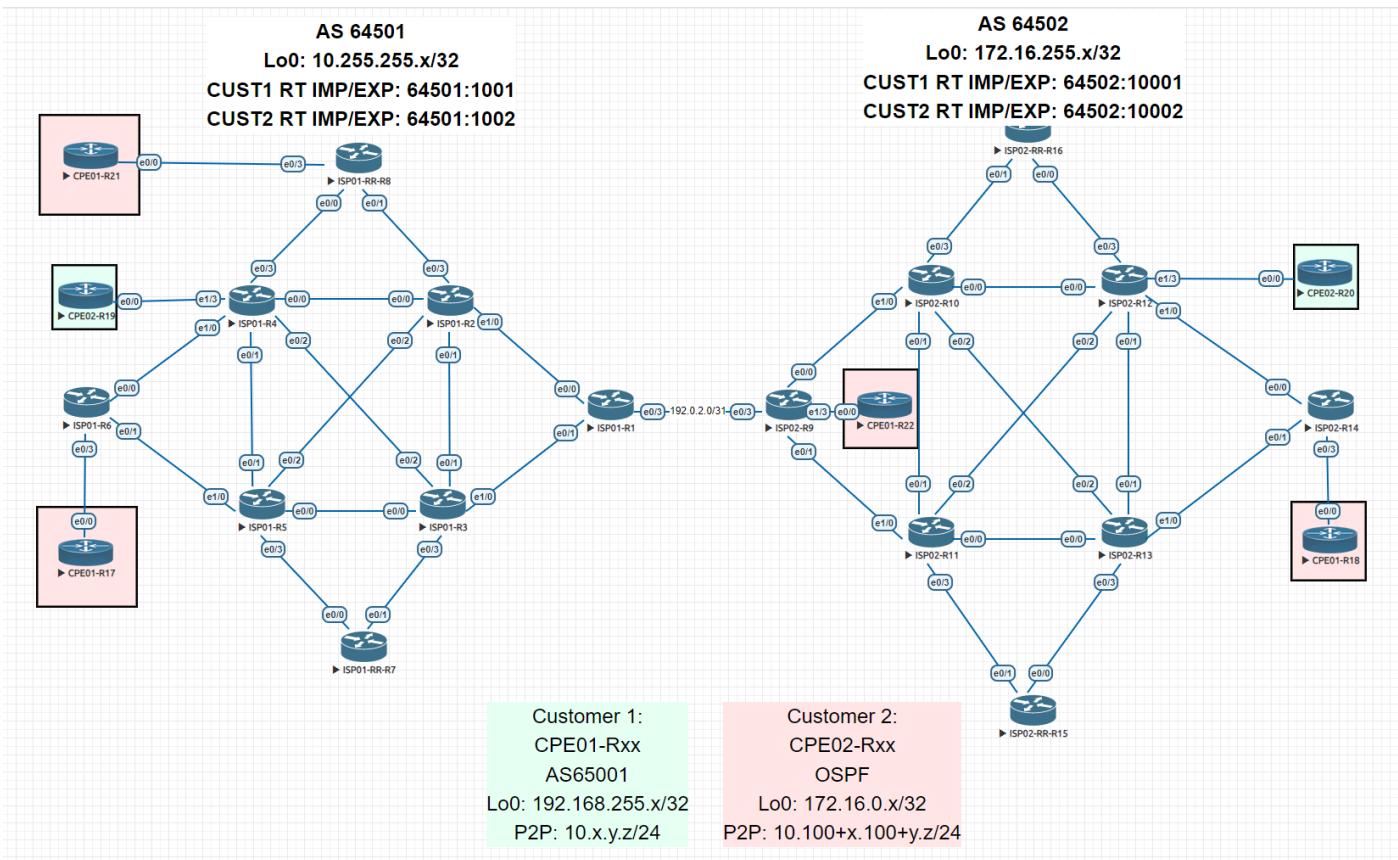


Cisco eBGP option C – MPLS L3VPN



Использованные образы:

i86bi_LinuxL3-AdvEnterpriseK9-M2_157_3_May_2018.bin

0. Summary

В данной лабораторной работе необходимо обеспечить связность между сайтами клиентов, подключенных к разным автономным системам. Для этого необходимо организовать interAS-стык по схеме «option C».

Данную задачу можно разбить на две последовательные подзадачи:

1. Поднятие interAS-стыка в AFI/SAFI IPv4 Labeled-unicast между ASBR'ами.
2. Поднятие interAS-стыков в AFI/SAFI VPNv4 Unicast между рефлекторами, расположенных в разных AS.

1. IPv4 Labeled Unicast

В первую очередь, нам необходимо организовать транспортные LSP между маршрутизаторами, расположеными в разных AS.

Для этого сначала поднимем в каждой AS iBGP-сессии в ipv4 labeled-unicast:

ISP01-ASBR-R1	ISP01-R2	ISP01-R3
<pre>router bgp 64501 address-family ipv4 network 10.255.255.1 mask 255.255.255.255 neighbor 10.255.255.7 activate neighbor 10.255.255.8 activate neighbor RR send-community neighbor RR send-label explicit- null exit-address-family ! !</pre>	<pre>router bgp 64501 address-family ipv4 network 10.255.255.2 mask 255.255.255.255 neighbor 10.255.255.7 activate neighbor 10.255.255.8 activate neighbor RR send-community neighbor RR send-label explicit- null exit-address-family ! !</pre>	<pre>router bgp 64501 address-family ipv4 network 10.255.255.3 mask 255.255.255.255 neighbor 10.255.255.7 activate neighbor 10.255.255.8 activate neighbor RR send-community neighbor RR send-label explicit- null exit-address-family ! !</pre>
ISP01-R4	ISP01-R5	ISP01-R6
<pre>router bgp 64501 address-family ipv4 network 10.255.255.4 mask 255.255.255.255 neighbor 10.255.255.7 activate neighbor 10.255.255.8 activate neighbor RR send-community neighbor RR send-label explicit- null exit-address-family ! !</pre>	<pre>router bgp 64501 address-family ipv4 network 10.255.255.5 mask 255.255.255.255 neighbor 10.255.255.7 activate neighbor 10.255.255.8 activate neighbor RR send-community neighbor RR send-label explicit- null exit-address-family ! !</pre>	<pre>router bgp 64501 address-family ipv4 network 10.255.255.6 mask 255.255.255.255 neighbor 10.255.255.7 activate neighbor 10.255.255.8 activate neighbor RR send-community neighbor RR send-label explicit- null exit-address-family ! !</pre>

ISP-RR-R7	ISP-RR-R8
<pre>router bgp 64501 ! address-family ipv4 network 10.255.255.7 mask 255.255.255.255 neighbor 10.255.255.1 activate neighbor 10.255.255.2 activate neighbor 10.255.255.3 activate neighbor 10.255.255.4 activate neighbor 10.255.255.5 activate neighbor 10.255.255.6 activate neighbor 10.255.255.8 activate neighbor OWN_RRC send-community neighbor OWN_RRC send-label explicit-null neighbor OWN_RRC route-reflector-client neighbor RRN send-community neighbor RRN send-label explicit-null exit-address-family ! !</pre>	<pre>router bgp 64501 ! address-family ipv4 network 10.255.255.8 mask 255.255.255.255 neighbor 10.255.255.1 activate neighbor 10.255.255.2 activate neighbor 10.255.255.3 activate neighbor 10.255.255.4 activate neighbor 10.255.255.5 activate neighbor 10.255.255.6 activate neighbor 10.255.255.7 activate neighbor OWN_RRC send-community neighbor OWN_RRC send-label explicit-null neighbor OWN_RRC route-reflector-client neighbor RRN send-community neighbor RRN send-label explicit-null exit-address-family ! !</pre>

Проверяем в AS64501:

```
ISP01-RR-R8#show bgp ipv4 unicast | b Network
      Network      Next Hop          Metric LocPrf Weight Path
r i 10.255.255.1/32 10.255.255.1      0    100    0 i
r>i                                10.255.255.1      0    100    0 i
r i 10.255.255.2/32 10.255.255.2      0    100    0 i
r>i                                10.255.255.2      0    100    0 i
r i 10.255.255.3/32 10.255.255.3      0    100    0 i
r>i                                10.255.255.3      0    100    0 i
r i 10.255.255.4/32 10.255.255.4      0    100    0 i
r>i                                10.255.255.4      0    100    0 i
r i 10.255.255.5/32 10.255.255.5      0    100    0 i
r>i                                10.255.255.5      0    100    0 i
r i 10.255.255.6/32 10.255.255.6      0    100    0 i
r>i                                10.255.255.6      0    100    0 i
r>i 10.255.255.7/32 10.255.255.7      0    100    0 i
*-> 10.255.255.8/32 0.0.0.0          0     0       32768 i
```

```
ISP01-RR-R8#show bgp ipv4 unicast labels
Network      Next Hop      In label/Out label
10.255.255.1/32 10.255.255.1    nolabel/exp-null
10.255.255.1/32 10.255.255.1    nolabel/exp-null
10.255.255.2/32 10.255.255.2    nolabel/exp-null
10.255.255.2/32 10.255.255.2    nolabel/exp-null
10.255.255.3/32 10.255.255.3    nolabel/exp-null
10.255.255.3/32 10.255.255.3    nolabel/exp-null
10.255.255.4/32 10.255.255.4    nolabel/exp-null
10.255.255.4/32 10.255.255.4    nolabel/exp-null
10.255.255.5/32 10.255.255.5    nolabel/exp-null
10.255.255.5/32 10.255.255.5    nolabel/exp-null
10.255.255.6/32 10.255.255.6    nolabel/exp-null
10.255.255.6/32 10.255.255.6    nolabel/exp-null
10.255.255.7/32 10.255.255.7    nolabel/exp-null
10.255.255.8/32 0.0.0.0      imp-null/nolabel
```

AS64502

ISP02-ASBR-R09	ISP02-R10	ISP02-R11
<pre>router bgp 64502 address-family ipv4 network 172.16.255.9 mask 255.255.255.255 neighbor 172.16.255.15 activate neighbor 172.16.255.16 activate neighbor RR send-community neighbor RR send-label explicit- null exit-address-family ! !</pre>	<pre>router bgp 64502 address-family ipv4 network 172.16.255.10 mask 255.255.255.255 neighbor 172.16.255.15 activate neighbor 172.16.255.16 activate neighbor RR send-community neighbor RR send-label explicit- null exit-address-family ! !</pre>	<pre>router bgp 64502 address-family ipv4 network 172.16.255.11 mask 255.255.255.255 neighbor 172.16.255.15 activate neighbor 172.16.255.16 activate neighbor RR send-community neighbor RR send-label explicit- null exit-address-family ! !</pre>
ISP02-R12	ISP02-R13	ISP02-R14
<pre>router bgp 64502 address-family ipv4 network 172.16.255.12 mask 255.255.255.255 neighbor 172.16.255.15 activate neighbor 172.16.255.16 activate neighbor RR send-community neighbor RR send-label explicit- null exit-address-family ! !</pre>	<pre>router bgp 64502 address-family ipv4 network 172.16.255.13 mask 255.255.255.255 neighbor 172.16.255.15 activate neighbor 172.16.255.16 activate neighbor RR send-community neighbor RR send-label explicit- null exit-address-family ! !</pre>	<pre>router bgp 64502 address-family ipv4 network 172.16.255.14 mask 255.255.255.255 neighbor 172.16.255.15 activate neighbor 172.16.255.16 activate neighbor RR send-community neighbor RR send-label explicit- null exit-address-family ! !</pre>

ISP02-RR-R15	ISP02-RR-R16
<pre>router bgp 64502 ! address-family ipv4 network 172.16.255.15 mask 255.255.255.255 neighbor 172.16.255.9 activate neighbor 172.16.255.10 activate neighbor 172.16.255.11 activate neighbor 172.16.255.12 activate neighbor 172.16.255.13 activate neighbor 172.16.255.14 activate neighbor OWN_RRC send-community neighbor OWN_RRC send-label explicit-null neighbor OWN_RRC route-reflector-client exit-address-family ! !</pre>	<pre>router bgp 64502 ! address-family ipv4 network 172.16.255.16 mask 255.255.255.255 neighbor 172.16.255.9 activate neighbor 172.16.255.10 activate neighbor 172.16.255.11 activate neighbor 172.16.255.12 activate neighbor 172.16.255.13 activate neighbor 172.16.255.14 activate neighbor OWN_RRC send-community neighbor OWN_RRC send-label explicit-null neighbor OWN_RRC route-reflector-client exit-address-family ! !</pre>

Проверяем в AS64502:

```
ISP02-RR-R15#show bgp ipv4 unicast | b Network
  Network      Next Hop      Metric LocPrf Weight Path
r>i 172.16.255.9/32 172.16.255.9          0    100    0 i
r>i 172.16.255.10/32 172.16.255.10         0    100    0 i
r>i 172.16.255.11/32 172.16.255.11         0    100    0 i
r>i 172.16.255.12/32 172.16.255.12         0    100    0 i
r>i 172.16.255.13/32 172.16.255.13         0    100    0 i
r>i 172.16.255.14/32 172.16.255.14         0    100    0 i
*> 172.16.255.15/32 0.0.0.0          0            32768 i

ISP02-RR-R15#show bgp ipv4 unicast labels
  Network      Next Hop      In label/Out label
172.16.255.9/32 172.16.255.9    nolabel/exp-null
```

```

172.16.255.10/32 172.16.255.10 nolabel/exp-null
172.16.255.11/32 172.16.255.11 nolabel/exp-null
172.16.255.12/32 172.16.255.12 nolabel/exp-null
172.16.255.13/32 172.16.255.13 nolabel/exp-null
172.16.255.14/32 172.16.255.14 nolabel/exp-null
172.16.255.15/32 0.0.0.0      imp-null/nolabel

```

Сейчас поднимем BGP-стык между ASBR'ми:

ISP01-ASBR-R1	ISP02-ASBR-R9
<pre> router bgp 64501 ! neighbor 192.0.2.1 remote-as 64502 neighbor 192.0.2.1 update-source Ethernet0/3 ! address-family ipv4 neighbor 192.0.2.1 activate neighbor 192.0.2.1 send-community neighbor 192.0.2.1 send-label explicit-null neighbor RR next-hop-self ! !</pre>	<pre> router bgp 64502 ! neighbor 192.0.2.0 remote-as 64501 neighbor 192.0.2.0 update-source Ethernet0/3 ! address-family ipv4 neighbor 192.0.2.0 activate neighbor 192.0.2.0 send-community neighbor 192.0.2.0 send-label explicit-null neighbor RR next-hop-self ! !</pre>

Как и в случае с optB, при поднятии BGP-сессии, подразумевающей обмен метками и прохождение MPLS-пакетов, включается MPLS на интерфейсе:

```

ISP01-ASBR-R1#
*Mar 19 13:18:45.522: %BGP-5-ADJCHANGE: neighbor 192.0.2.1 Up
ISP01-ASBR-R1#
*Mar 19 13:18:45.525: %BGP_LMM-6-AUTOGEN1: The mpls bgp forwarding command has been configured on interface:
Ethernet0/3
ISP01-ASBR-R1#show mpls interface
Interface          IP           Tunnel   BGP  Static Operational
Ethernet0/0        Yes (ldp)    No       No   No    Yes
Ethernet0/1        Yes (ldp)    No       No   No    Yes
Ethernet0/3        No          No       Yes  No    Yes
ISP01-ASBR-R1#show ip route 192.0.2.0
Routing entry for 192.0.2.0/24, 3 known subnets
Attached (3 connections)
  Variably subnetted with 2 masks
C      192.0.2.0/31 is directly connected, Ethernet0/3
L      192.0.2.0/32 is directly connected, Ethernet0/3
C      192.0.2.1/32 is directly connected, Ethernet0/3
ISP01-ASBR-R1#

```

Проверяем RIB и LFIB на ASBR:

```

ISP01-ASBR-R1#show ip route bgp | b Gateway
Gateway of last resort is not set

  172.16.0.0/32 is subnetted, 8 subnets
B     172.16.255.9 [20/0] via 192.0.2.1, 00:06:26
B     172.16.255.10 [20/0] via 192.0.2.1, 00:06:26
B     172.16.255.11 [20/0] via 192.0.2.1, 00:06:26
B     172.16.255.12 [20/0] via 192.0.2.1, 00:06:26
B     172.16.255.13 [20/0] via 192.0.2.1, 00:06:26
B     172.16.255.14 [20/0] via 192.0.2.1, 00:06:26
B     172.16.255.15 [20/0] via 192.0.2.1, 00:06:26
B     172.16.255.16 [20/0] via 192.0.2.1, 00:06:26
ISP01-ASBR-R1#show mpls forwarding-table next-hop 192.0.2.1
Local    Outgoing   Prefix      Bytes Label  Outgoing      Next Hop
Label    Label     or Tunnel Id   Switched   interface
10000    Pop Label 192.0.2.1/32  0          Et0/3      192.0.2.1
10001    explicit-n 172.16.255.9/32 0          Et0/3      192.0.2.1
10002    90009     172.16.255.10/32 0          Et0/3      192.0.2.1
10011    90008     172.16.255.11/32 0          Et0/3      192.0.2.1
10012    90007     172.16.255.12/32 0          Et0/3      192.0.2.1
10013    90006     172.16.255.13/32 0          Et0/3      192.0.2.1
10014    90005     172.16.255.14/32 0          Et0/3      192.0.2.1
10015    90004     172.16.255.15/32 0          Et0/3      192.0.2.1
10016    90003     172.16.255.16/32 0          Et0/3      192.0.2.1
ISP01-ASBR-R1#

```

Проверяем связность:

```
ISP01-ASBR-R1#ping 172.16.255.9 source 10.255.255.1
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 172.16.255.9, timeout is 2 seconds:
Packet sent with a source address of 10.255.255.1
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 1/1/1 ms
ISP01-ASBR-R1#ping 172.16.255.15 source 10.255.255.1
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 172.16.255.15, timeout is 2 seconds:
Packet sent with a source address of 10.255.255.1
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 1/1/2 ms
```

Теперь посмотрим тоже самое на RR:

```
ISP01-RR-R8#show ip route bgp | b Gateway
Gateway of last resort is not set

  172.16.0.0/32 is subnetted, 8 subnets
B    172.16.255.9 [200/0] via 10.255.255.1, 00:16:01
B    172.16.255.10 [200/0] via 10.255.255.1, 00:16:01
B    172.16.255.11 [200/0] via 10.255.255.1, 00:16:01
B    172.16.255.12 [200/0] via 10.255.255.1, 00:16:01
B    172.16.255.13 [200/0] via 10.255.255.1, 00:16:01
B    172.16.255.14 [200/0] via 10.255.255.1, 00:16:01
B    172.16.255.15 [200/0] via 10.255.255.1, 00:16:01
B    172.16.255.16 [200/0] via 10.255.255.1, 00:16:01
ISP01-RR-R8#show mpls forwarding-table next-hop 10.255.255.1
Local      Outgoing   Prefix          Bytes Label  Outgoing      Next Hop
Label      Label      or Tunnel Id  Switched     interface
ISP01-RR-R8#show mpls forwarding-table
Local      Outgoing   Prefix          Bytes Label  Outgoing      Next Hop
Label      Label      or Tunnel Id  Switched     interface
80002    explicit-n  10.255.255.4/32 0        Et0/0       10.4.8.4
80003    40008      10.255.255.6/32 0        Et0/0       10.4.8.4
80004    40006      10.255.255.5/32 0        Et0/0       10.4.8.4
80005    20009      10.255.255.7/32 0        Et0/1       10.2.8.2
        40011      10.255.255.7/32 0        Et0/0       10.4.8.4
80006    explicit-n  10.255.255.2/32 0        Et0/1       10.2.8.2
80007    20005      10.255.255.1/32 0        Et0/1       10.2.8.2
80008    20006      10.255.255.3/32 0        Et0/1       10.2.8.2
80009    Pop Label   10.0.0.8/32[V] 0        aggregate/MGMT
80010    No Label    10.8.21.0/24[V] 0        aggregate/CUST1
80011    No Label    192.168.255.21/32[V] \
                                         0        Et0/3       10.8.21.21
ISP01-RR-R8#
```

В RIB маршруты есть, а вот в LFIB, как ни странно, соответствия меток и префиксов нужных нет.
Посмотрим, что у нас видно в BGP:

```
ISP01-RR-R8#show bgp ipv4 unicast labels
Network      Next Hop      In label/Out label
10.255.255.1/32 10.255.255.1  nolabel/exp-null
                10.255.255.1  nolabel/exp-null
10.255.255.2/32 10.255.255.2  nolabel/exp-null
                10.255.255.2  nolabel/exp-null
10.255.255.3/32 10.255.255.3  nolabel/exp-null
                10.255.255.3  nolabel/exp-null
10.255.255.4/32 10.255.255.4  nolabel/exp-null
                10.255.255.4  nolabel/exp-null
10.255.255.5/32 10.255.255.5  nolabel/exp-null
                10.255.255.5  nolabel/exp-null
10.255.255.6/32 10.255.255.6  nolabel/exp-null
                10.255.255.6  nolabel/exp-null
10.255.255.7/32 10.255.255.7  nolabel/exp-null
10.255.255.8/32 0.0.0.0      imp-null/nolabel
172.16.255.9/32 10.255.255.1  nolabel/10001
                10.255.255.1  nolabel/10001
172.16.255.10/32 10.255.255.1  nolabel/10002
                10.255.255.1  nolabel/10002
172.16.255.11/32 10.255.255.1  nolabel/10011
                10.255.255.1  nolabel/10011
172.16.255.12/32 10.255.255.1  nolabel/10012
                10.255.255.1  nolabel/10012
172.16.255.13/32 10.255.255.1  nolabel/10013
                10.255.255.1  nolabel/10013
```

```

172.16.255.14/32 10.255.255.1      nolabel/10014
          10.255.255.1      nolabel/10014
172.16.255.15/32 10.255.255.1      nolabel/10015
          10.255.255.1      nolabel/10015
172.16.255.16/32 10.255.255.1      nolabel/10016
          10.255.255.1      nolabel/10016

```

Здесь мы видим, что метки для лупбэков из AS64502 получены. Посмотрим более детально на какой-нибудь из префиксов:

```

ISP01-RR-R8#show bgp ipv4 unicast 172.16.255.9/32
BGP routing table entry for 172.16.255.9/32, version 10
Paths: (2 available, best #2, table default)
  Advertised to update-groups:
    1             2
  Refresh Epoch 1
  64502
    10.255.255.1 (metric 200) from 10.255.255.7 (10.255.255.7)
      Origin IGP, metric 0, localpref 100, valid, internal
      Originator: 10.255.255.1, Cluster list: 10.255.255.7
      mpls labels in/out nolabel/10001
        rx pathid: 0, tx pathid: 0
  Refresh Epoch 5
  64502, (Received from a RR-client)
    10.255.255.1 (metric 200) from 10.255.255.1 (10.255.255.1)
      Origin IGP, metric 0, localpref 100, valid, internal, best
      mpls labels in/out nolabel/10001
        rx pathid: 0, tx pathid: 0x0
ISP01-RR-R8#

```

Мы видим, что маршрут валиден, нормально резолвится через 10.255.255.1.

Давайте глянем на то, что видно в FIB:

```

ISP01-RR-R8#show ip cef 172.16.255.9/32 detail
172.16.255.9/32, epoch 0, flags [rib defined all labels]
  recursive via 10.255.255.1 label 10001
    nexthop 10.2.8.2 Ethernet0/1 label 20005-(local:80007)
ISP01-RR-R8#

```

Здесь мы видим, что используется стек меток: внешняя метка - это метка next-hop'a 10.255.255.1, а внутренняя - это метка, полученная по BGP вместе с префиксом. Т.о., в LSP R8 > R16 используется пара меток на участке R8 - R1, а после выхода с R1 остается только одна метка.

Проверим связность:

```

ISP01-RR-R8#ping 172.16.255.16 source Lo0
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 172.16.255.16, timeout is 2 seconds:
Packet sent with a source address of 10.255.255.8
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 2/2/2 ms
ISP01-RR-R8#traceroute 172.16.255.16 source Lo0
Type escape sequence to abort.
Tracing the route to 172.16.255.16
VRF info: (vrf in name/id, vrf out name/id)
  1 10.2.8.2 [MPLS: Labels 20005/10016 Exp 0] 3 msec 2 msec 3 msec
  2 10.1.2.1 [MPLS: Labels 0/10016 Exp 0] 3 msec 2 msec 2 msec
  3 192.0.2.1 [MPLS: Label 90003 Exp 0] 3 msec 3 msec 3 msec
  4 172.31.109.10 [MPLS: Label 100005 Exp 0] 3 msec 2 msec 3 msec
  5 172.31.116.16 2 msec * 3 msec
ISP01-RR-R8#

```

Собственно, трейс подтверждает вышесказанное. На R1 происходит своп метки 10016 на 90003 и отдается в сторону R9:

```

ISP01-ASBR-R1#show mpls forwarding-table labels 10016
Local      Outgoing   Prefix           Bytes Label  Outgoing      Next Hop
Label      Label       or Tunnel Id   Switched     interface
10016      90003      172.16.255.16/32 2380          Et0/3       192.0.2.1
ISP01-ASBR-R1#

```

2. VPNv4 Unicast RR <> RR

На данный момент у нас есть связность между рефлекторами, поэтому приступаем к настройке eBGP-сессий в VPNv4 Unicast.

ISP01-RR-R7 / ISP01-RR-R8	ISP02-RR-R15 / ISP02-RR-R16
<pre> router bgp 64501 neighbor RR_OPT peer-group neighbor RR_OPT remote-as 64502 neighbor RR_OPT update-source Loopback0 neighbor RR_OPT ebgp-multipath 255 neighbor 172.16.255.15 peer-group RR_OPT neighbor 172.16.255.16 peer-group RR_OPT ! address-family vpnv4 unicast neighbor 172.16.255.15 activate neighbor 172.16.255.16 activate neighbor RR_OPT send-community both ! !</pre>	<pre> router bgp 64502 neighbor RR_OPT peer-group neighbor RR_OPT remote-as 64501 neighbor RR_OPT update-source Loopback0 neighbor RR_OPT ebgp-multipath 255 neighbor 10.255.255.7 peer-group RR_OPT neighbor 10.255.255.8 peer-group RR_OPT ! address-family vpnv4 unicast neighbor 10.255.255.7 activate neighbor 10.255.255.8 activate neighbor RR_OPT send-community both ! !</pre>

Проверим на R8:

```

ISP01-RR-R8#show bgp vpnv4 unicast all summary | i 64502
172.16.255.15    4      64502      16      17      142      0      0 00:00:47      13
172.16.255.16    4      64502      17      19      142      0      0 00:00:35      13
ISP01-RR-R8#show bgp vpnv4 unicast all neighbors 172.16.255.15 routes | b Netw
      Network          Next Hop          Metric LocPrf Weight Path
Route Distinguisher: 172.16.255.9:11
  *> 10.0.0.9/32      172.16.255.15      0 64502 ?
Route Distinguisher: 172.16.255.9:10001
  *> 10.9.22.0/24     172.16.255.15      0 64502 ?
  *> 192.168.255.22/32
                           172.16.255.15      0 64502 65001 i
Route Distinguisher: 172.16.255.10:11
  *> 10.0.0.10/32     172.16.255.15      0 64502 ?
Route Distinguisher: 172.16.255.11:11
  *> 10.0.0.11/32     172.16.255.15      0 64502 ?
Route Distinguisher: 172.16.255.12:11
  *> 10.0.0.12/32     172.16.255.15      0 64502 ?
Route Distinguisher: 172.16.255.12:10002
  *> 10.112.120.0/24   172.16.255.15      0 64502 ?
  *> 172.16.0.20/32    172.16.255.15      0 64502 ?
Route Distinguisher: 172.16.255.13:11
  *> 10.0.0.13/32     172.16.255.15      0 64502 ?
Route Distinguisher: 172.16.255.14:11
  *> 10.0.0.14/32     172.16.255.15      0 64502 ?
Route Distinguisher: 172.16.255.14:10001
  *> 10.14.18.0/24     172.16.255.15      0 64502 ?
  *> 192.168.255.18/32
                           172.16.255.15      0 64502 65001 i
Route Distinguisher: 172.16.255.15:11
  *> 10.0.0.15/32     172.16.255.15      0      0 64502 ?

Total number of prefixes 13
ISP01-RR-R8#show bgp vpnv4 unicast all 192.168.255.18/32
BGP routing table entry for 172.16.255.14:10001:192.168.255.18/32, version 140
Paths: (3 available, best #3, no table)
  Advertised to update-groups:
    6      7      8
  Refresh Epoch 2
  64502 65001
    172.16.255.16 (metric 200) (via default) from 172.16.255.16 (172.16.255.16)
      Origin IGP, localpref 100, valid, external
      Extended Community: RT:64502:10001
      mpls labels in/out 80030/160001
      rx pathid: 0, tx pathid: 0
  Refresh Epoch 1
  64502 65001
    172.16.255.16 (metric 200) (via default) from 10.255.255.7 (10.255.255.7)
      Origin IGP, metric 0, localpref 100, valid, internal
      Extended Community: RT:64502:10001
      mpls labels in/out 80030/160001
      rx pathid: 0, tx pathid: 0
  Refresh Epoch 1
  64502 65001
    172.16.255.15 (metric 200) (via default) from 172.16.255.15 (172.16.255.15)
      Origin IGP, localpref 100, valid, external, best

```

```

Extended Community: RT:64502:10001
mpls labels in/out 80030/150019
rx pathid: 0, tx pathid: 0x0
ISP01-RR-R8#

```

Control-plane у нас сошелся, проверим теперь работу data plane. Для этого сначала для клиентских vrf добавим соответствующие RT из AS64502.

ISP01-R4	ISP01-R6	ISP01-RR-R8
vrf definition CUST2 ! address-family ipv4 route-target export 64502:10002 route-target import 64502:10002 exit-address-family !	vrf definition CUST1 ! address-family ipv4 route-target export 64502:10001 route-target import 64502:10001 exit-address-family !	vrf definition CUST1 ! address-family ipv4 route-target export 64502:10001 route-target import 64502:10001 exit-address-family !

Для начала посмотрим что у нас есть на PE на примере vrf CUST2, на R4:

```

ISP01-R4#show ip route vrf CUST2 | b Gateway
Gateway of last resort is not set

      10.0.0.0/8 is variably subnetted, 3 subnets, 2 masks
C          10.104.119.0/24 is directly connected, Ethernet1/3
L          10.104.119.104/32 is directly connected, Ethernet1/3
B          10.112.120.0/24 [200/0] via 172.16.255.16, 00:14:55
              172.16.0.0/32 is subnetted, 2 subnets
O            172.16.0.19 [110/11] via 10.104.119.19, 03:52:04, Ethernet1/3
B            172.16.0.20 [200/0] via 172.16.255.16, 00:14:55
ISP01-R4#
ISP01-R4#show bgp vpnv4 unicast vrf CUST2 | b Network
      Network          Next Hop          Metric LocPrf Weight Path
Route Distinguisher: 10.255.255.4:1002 (default for vrf CUST2)
*-> 10.104.119.0/24  0.0.0.0          0        32768 ?
*->i 10.112.120.0/24 172.16.255.16    0        100      0 64502 ?
*-> 172.16.0.19/32   10.104.119.19    11       32768 ?
*->i 172.16.0.20/32   172.16.255.16    0        100      0 64502 ?
ISP01-R4#
ISP01-R4#show bgp vpnv4 unicast vrf CUST2 172.16.0.20/32
BGP routing table entry for 10.255.255.4:1002:172.16.0.20/32, version 105
Paths: (1 available, best #1, table CUST2)
Flag: 0x100
Not advertised to any peer
Refresh Epoch 3
64502, imported path from 172.16.255.12:10002:172.16.0.20/32 (global)
    172.16.255.16 (metric 200) (via default) from 10.255.255.7 (10.255.255.7)
        Origin incomplete, metric 0, localpref 100, valid, internal, best
        Extended Community: RT:64502:10002 OSPF DOMAIN ID:0x0005:0x00027120200
            OSPF RT:0.0.0.0:2:0 OSPF ROUTER ID:10.112.120.12:0
            mpls labels in/out nolabel/160011
            rx pathid: 0, tx pathid: 0x0
ISP01-R4#

```

Здесь мы видим, что маршруты были получены от наших RR, next-hop'ом считается 172.16.255.16, VPN-метка 160011.

Давайте попробуем проверить связность между CPE02-R19 и CPE02-R20:

```

CPE02-R19#ping 172.16.0.20 source Lo0
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 172.16.0.20, timeout is 2 seconds:
Packet sent with a source address of 172.16.0.19
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 4/4/5 ms
CPE02-R19#traceroute 172.16.0.20 source Lo0
Type escape sequence to abort.
Tracing the route to 172.16.0.20
VRF info: (vrf in name/id, vrf out name/id)
 1 10.104.119.104 1 msec 0 msec 1 msec
 2 10.2.4.2 [MPLS: Labels 20005/10016/160011 Exp 0] 5 msec 4 msec 4 msec
 3 10.1.2.1 [MPLS: Labels 0/10016/160011 Exp 0] 3 msec 3 msec 4 msec
 4 192.0.2.1 [MPLS: Labels 90003/160011 Exp 0] 4 msec 4 msec 4 msec
 5 172.31.109.10 [MPLS: Labels 100005/160011 Exp 0] 4 msec 3 msec 3 msec
 6 172.31.116.16 [MPLS: Labels 0/160011 Exp 0] 4 msec 4 msec 3 msec
 7 10.112.120.12 [MPLS: Labels 0/120001 Exp 0] 3 msec 2 msec 3 msec
 8 10.112.120.20 3 msec * 4 msec

```

Связность есть. Давайте рассмотрим трейс поподробней:

№	Запись	Комментарий
1	10.104.119.104	Пакет от CPE02-R19 до ISP01-R4
2	10.2.4.2 [MPLS: Labels 20005/10016/160011 Exp 0]	Пакет от ISP01-R4 до ISP01-R2: PUSH на R4 метку 160011 – VPN-метка префикса 172.16.0.20/32 PUSH на R4 метку 10016 – метка префикса 172.16.255.16/32 PUSH на R4 метку 20005 – метка префикса 10.255.255.1
3	10.1.2.1 [MPLS: Labels 0/10016/160011 Exp 0]	Пакет от ISP01-R2 до ISP01-ASBR-R1: SWAP на R2 метки 20005 на 0 (exp-null)
4	192.0.2.1 [MPLS: Labels 90003/160011 Exp 0]	Пакет от ISP01-ASBR-R1 до ISP02-ASBR-R9: POP на R1 метки 0 SWAP метки 10016 на 90003
5	172.31.109.10 [MPLS: Labels 100005/160011 Exp 0]	Пакет от ISP02-ASBR-R9 до ISP02-R10: SWAP на R9 метки 90003 на 100005
6	172.31.116.16 [MPLS: Labels 0/160011 Exp 0]	Пакет от ISP02-R10 до ISP02-RR-R16: SWAP на R10 метки 100005 на 0 (exp-null)
7	10.112.120.12 [MPLS: Labels 0/120001 Exp 0]	Пакет от ISP02-RR-R16 до ISP02-R12: POP на R16 метки 0 SWAP на R16 метки 160011 на 120001 PUSH на R16 метки 0 (exp-null)
8	10.112.120.20	Пакет от ISP02-R12 до CPE02-R20: POP на R12 метки 0 POP на R12 метки 120001

Собственно, из трейса мы видим, что трафик до R20 проходит через ISP02-RR-R16. Визуально, конечно, это выглядит субоптимальным путем, однако, вполне возможно, что PE R12 действительно доступен через RR. На всякий случай сделаем трейс от R4 до R12:

```
ISP01-R4#traceroute 172.16.255.12 source Loopback0
Type escape sequence to abort.
Tracing the route to 172.16.255.12
VRF info: (vrf in name/id, vrf out name/id)
 1 10.2.4.2 [MPLS: Labels 20005/10012 Exp 0] 2 msec 2 msec 1 msec
 2 10.1.2.1 [MPLS: Labels 0/10012 Exp 0] 2 msec 2 msec 2 msec
 3 192.0.2.1 [MPLS: Label 90007 Exp 0] 2 msec 2 msec 2 msec
 4 172.31.109.10 [MPLS: Label 100009 Exp 0] 2 msec 2 msec 2 msec
 5 172.31.112.12 3 msec * 3 msec
ISP01-R4#
```

По трейсу видно, что до ISP02-R12 трафик проходит напрямую от R10, не проходя через ISP02-RR-R16.

Давайте разбираться, почему так получается. Сначала посмотрим на R4:

```
ISP01-R4#show bgp vpnv4 unicast vrf CUST2 172.16.0.20/32
BGP routing table entry for 10.255.255.4:1002:172.16.0.20/32, version 105
Paths: (1 available, best #1, table CUST2)
Flag: 0x100
    Not advertised to any peer
Refresh Epoch 3
64502, imported path from 172.16.255.12:10002:172.16.0.20/32 (global)
    172.16.255.16 (metric 200) (via default) from 10.255.255.7 (10.255.255.7)
        Origin incomplete, metric 0, localpref 100, valid, internal, best
        Extended Community: RT:64502:10002 OSPF DOMAIN ID:0x0005:0x000027120200
            OSPF RT:0.0.0.0:2:0 OSPF ROUTER ID:10.112.120.12:0
            mpls labels in/out nolabel/160011
            rx pathid: 0, tx pathid: 0x0
ISP01-R4#
```

Здесь мы видим, что маршрут мы получили от нашего RR 10.255.255.7, в качестве next-hop стоит 172.16.255.16. Почему? А потому что у нас созданы eBGP VPNv4 сессии между RR. Т.к. это eBGP, то по умолчанию в качестве next-hop проставляется адрес соседа, от которого эти маршруты получены (в нашем случае это 172.16.255.15 и 172.16.255.16).

В общем-то, на данный момент все работает, но, все же, лучше привести к нормальному виду. Для этого нам необходимо на RR включить команду next-hop-unchanged в сторону RR из соседней AS (ну и не забыть переанонсировать маршруты!).

ISP01-RR-R7 / ISP01-RR-R8	ISP02-RR-15 / ISP02-RR-16
router bgp 64501 ! address-family vpnv4 neighbor RR_OPT next-hop-unchanged	router bgp 64501 ! address-family vpnv4 neighbor RR_OPT next-hop-unchanged

exit-address-family ! clear bgp vpng4 unicast * out	exit-address-family ! clear bgp vpng4 unicast * out
---	---

После того, как мы выполнили эти настройки, проверим, что сейчас видно на R4:

```
ISP01-R4#show bgp vpng4 unicast vrf CUST2 | b Network
   Network          Next Hop          Metric LocPrf Weight Path
Route Distinguisher: 10.255.255.4:1002 (default for vrf CUST2)
  *> 10.104.119.0/24  0.0.0.0          0      32768 ?
  *>i 10.112.120.0/24 172.16.255.12    0     100      0 64502 ?
  *> 172.16.0.19/32  10.104.119.19    11      32768 ?
  *>i 172.16.0.20/32 172.16.255.12    0     100      0 64502 ?
ISP01-R4#
```

Сейчас у нас в качестве next-hop'ов стоит R12, чего мы и добивались. Проверим прохождение трафика еще и трейсом с CPE02-R19:

```
CPE02-R19#traceroute 172.16.0.20 source Lo0
Type escape sequence to abort.
Tracing the route to 172.16.0.20
VRF info: (vrf in name/id, vrf out name/id)
 1 10.104.119.104 2 msec 35 msec 1 msec
 2 10.2.4.2 [MPLS: Labels 20005/10012/120001 Exp 0] 8 msec 17 msec 18 msec
 3 10.1.2.1 [MPLS: Labels 0/10012/120001 Exp 0] 19 msec 8 msec 8 msec
 4 192.0.2.1 [MPLS: Labels 90007/120001 Exp 0] 8 msec 7 msec 5 msec
 5 172.31.109.10 [MPLS: Labels 100009/120001 Exp 0] 4 msec 3 msec 2 msec
 6 10.112.120.12 [MPLS: Labels 0/120001 Exp 0] 3 msec 2 msec 2 msec
 7 10.112.120.20 3 msec * 9 msec
CPE02-R19#
```

На данный момент трафик на R12 идет напрямую с R10, т.о. мы получили оптимальный путь прохождения трафика.

3. Option C solution config:

AS64501:

```
#ISP01-ASBR-R1:  
router bgp 64501  
neighbor 192.0.2.1 remote-as 64502  
neighbor 192.0.2.1 update-source Ethernet0/3  
!  
address-family ipv4  
network 10.255.255.1 mask 255.255.255.255  
neighbor 10.255.255.7 activate  
neighbor 10.255.255.8 activate  
neighbor 192.0.2.1 activate  
neighbor 192.0.2.1 send-community  
neighbor 192.0.2.1 send-label explicit-null  
neighbor RR next-hop-self  
neighbor RR send-community  
neighbor RR send-label explicit-null  
exit-address-family  
!  
!
```

#ISP01-R2:

```
router bgp 64501  
address-family ipv4  
network 10.255.255.2 mask 255.255.255.255  
neighbor 10.255.255.7 activate  
neighbor 10.255.255.8 activate  
neighbor RR send-community  
neighbor RR send-label explicit-null  
exit-address-family  
!  
!
```

#ISP01-R3:

```
router bgp 64501  
address-family ipv4  
network 10.255.255.3 mask 255.255.255.255  
neighbor 10.255.255.7 activate  
neighbor 10.255.255.8 activate  
neighbor RR send-community  
neighbor RR send-label explicit-null  
exit-address-family  
!  
!
```

#ISP01-R4:

```
vrf definition CUST2  
!  
address-family ipv4  
route-target export 64502:10002  
route-target import 64502:10002  
exit-address-family  
!
```

```
router bgp 64501  
address-family ipv4  
network 10.255.255.4 mask 255.255.255.255  
neighbor 10.255.255.7 activate  
neighbor 10.255.255.8 activate  
neighbor RR send-community  
neighbor RR send-label explicit-null  
exit-address-family  
!  
!
```

#ISP01-R5:

```
router bgp 64501  
address-family ipv4  
network 10.255.255.5 mask 255.255.255.255  
neighbor 10.255.255.7 activate  
neighbor 10.255.255.8 activate  
neighbor RR send-community  
neighbor RR send-label explicit-null  
exit-address-family  
!  
!
```

#ISP01-R6:

```
vrf definition CUST1  
!  
address-family ipv4
```

```

route-target export 64502:10001
route-target import 64502:10001
exit-address-family
!
router bgp 64501
address-family ipv4
network 10.255.255.6 mask 255.255.255.255
neighbor 10.255.255.7 activate
neighbor 10.255.255.8 activate
neighbor RR send-community
neighbor RR send-label explicit-null
exit-address-family
!
!
#ISP-RR-R7:
router bgp 64501
neighbor RR_OPT peer-group
neighbor RR_OPT remote-as 64502
neighbor RR_OPT update-source Loopback0
neighbor RR_OPT ebgp-multipath 255
neighbor 172.16.255.15 peer-group RR_OPT
neighbor 172.16.255.16 peer-group RR_OPT
!
address-family ipv4
network 10.255.255.7 mask 255.255.255.255
neighbor 10.255.255.1 activate
neighbor 10.255.255.2 activate
neighbor 10.255.255.3 activate
neighbor 10.255.255.4 activate
neighbor 10.255.255.5 activate
neighbor 10.255.255.6 activate
neighbor 10.255.255.8 activate
neighbor OWN_RRC send-community
neighbor OWN_RRC send-label explicit-null
neighbor OWN_RRC route-reflector-client
neighbor RRN send-community
neighbor RRN send-label explicit-null
exit-address-family
!
address-family vpng4 unicast
neighbor 172.16.255.15 activate
neighbor 172.16.255.16 activate
neighbor RR_OPT send-community both
neighbor RR_OPT next-hop-unchanged
!
!
#ISP-RR-R8:
vrf definition CUST1
!
address-family ipv4
route-target export 64502:10001
route-target import 64502:10001
exit-address-family
!
router bgp 64501
neighbor RR_OPT peer-group
neighbor RR_OPT remote-as 64502
neighbor RR_OPT update-source Loopback0
neighbor RR_OPT ebgp-multipath 255
neighbor 172.16.255.15 peer-group RR_OPT
neighbor 172.16.255.16 peer-group RR_OPT
!
address-family ipv4
network 10.255.255.8 mask 255.255.255.255
neighbor 10.255.255.1 activate
neighbor 10.255.255.2 activate
neighbor 10.255.255.3 activate
neighbor 10.255.255.4 activate
neighbor 10.255.255.5 activate
neighbor 10.255.255.6 activate
neighbor 10.255.255.7 activate
neighbor OWN_RRC send-community
neighbor OWN_RRC send-label explicit-null
neighbor OWN_RRC route-reflector-client
neighbor RRN send-community
neighbor RRN send-label explicit-null
exit-address-family
!
address-family vpng4 unicast
neighbor 172.16.255.15 activate

```

```
neighbor 172.16.255.16 activate
neighbor RR_OPT send-community both
neighbor RR_OPT next-hop-unchanged
!
```

AS64502:

```
#ISP02-ASBR-R09:
router bgp 64502
neighbor 192.0.2.0 remote-as 64501
neighbor 192.0.2.0 update-source Ethernet0/3
!
address-family ipv4
network 172.16.255.9 mask 255.255.255.255
neighbor 172.16.255.15 activate
neighbor 172.16.255.16 activate
neighbor 192.0.2.0 activate
neighbor 192.0.2.0 send-community
neighbor 192.0.2.0 send-label explicit-null
neighbor RR send-community
neighbor RR send-label explicit-null
neighbor RR next-hop-self
exit-address-family
!
```

```
!
```

#ISP02-R10:

```
router bgp 64502
address-family ipv4
network 172.16.255.10 mask 255.255.255.255
neighbor 172.16.255.15 activate
neighbor 172.16.255.16 activate
neighbor RR send-community
neighbor RR send-label explicit-null
exit-address-family
!
```

```
!
```

#ISP02-R11:

```
router bgp 64502
address-family ipv4
network 172.16.255.11 mask 255.255.255.255
neighbor 172.16.255.15 activate
neighbor 172.16.255.16 activate
neighbor RR send-community
neighbor RR send-label explicit-null
exit-address-family
!
```

```
!
```

#ISP02-R12:

```
router bgp 64502
address-family ipv4
network 172.16.255.12 mask 255.255.255.255
neighbor 172.16.255.15 activate
neighbor 172.16.255.16 activate
neighbor RR send-community
neighbor RR send-label explicit-null
exit-address-family
!
```

```
!
```

#ISP02-R13:

```
router bgp 64502
address-family ipv4
network 172.16.255.13 mask 255.255.255.255
neighbor 172.16.255.15 activate
neighbor 172.16.255.16 activate
neighbor RR send-community
neighbor RR send-label explicit-null
exit-address-family
!
```

```
!
```

#ISP02-R14:

```
router bgp 64502
address-family ipv4
network 172.16.255.14 mask 255.255.255.255
neighbor 172.16.255.15 activate
neighbor 172.16.255.16 activate
neighbor RR send-community
neighbor RR send-label explicit-null
```

```

exit-address-family
!
!

#ISP02-RR-R15:
router bgp 64502
neighbor RR_OPT peer-group
neighbor RR_OPT remote-as 64501
neighbor RR_OPT update-source Loopback0
neighbor RR_OPT ebgp-multipath 255
neighbor 10.255.255.7 peer-group RR_OPT
neighbor 10.255.255.8 peer-group RR_OPT
!
address-family ipv4
network 172.16.255.15 mask 255.255.255.255
neighbor 172.16.255.9 activate
neighbor 172.16.255.10 activate
neighbor 172.16.255.11 activate
neighbor 172.16.255.12 activate
neighbor 172.16.255.13 activate
neighbor 172.16.255.14 activate
neighbor OWN_RRC send-community
neighbor OWN_RRC send-label explicit-null
neighbor OWN_RRC route-reflector-client
exit-address-family
!
address-family vpng4 unicast
neighbor 10.255.255.7 activate
neighbor 10.255.255.8 activate
neighbor RR_OPT send-community both
neighbor RR_OPT next-hop-unchanged
!
!
#ISP02-RR-R16:
router bgp 64502
neighbor RR_OPT peer-group
neighbor RR_OPT remote-as 64501
neighbor RR_OPT update-source Loopback0
neighbor RR_OPT ebgp-multipath 255
neighbor 10.255.255.7 peer-group RR_OPT
neighbor 10.255.255.8 peer-group RR_OPT
!
address-family ipv4
network 172.16.255.16 mask 255.255.255.255
neighbor 172.16.255.9 activate
neighbor 172.16.255.10 activate
neighbor 172.16.255.11 activate
neighbor 172.16.255.12 activate
neighbor 172.16.255.13 activate
neighbor 172.16.255.14 activate
neighbor OWN_RRC send-community
neighbor OWN_RRC send-label explicit-null
neighbor OWN_RRC route-reflector-client
exit-address-family
!
address-family vpng4 unicast
neighbor 10.255.255.7 activate
neighbor 10.255.255.8 activate
neighbor RR_OPT send-community both
neighbor RR_OPT next-hop-unchanged
!
!
```