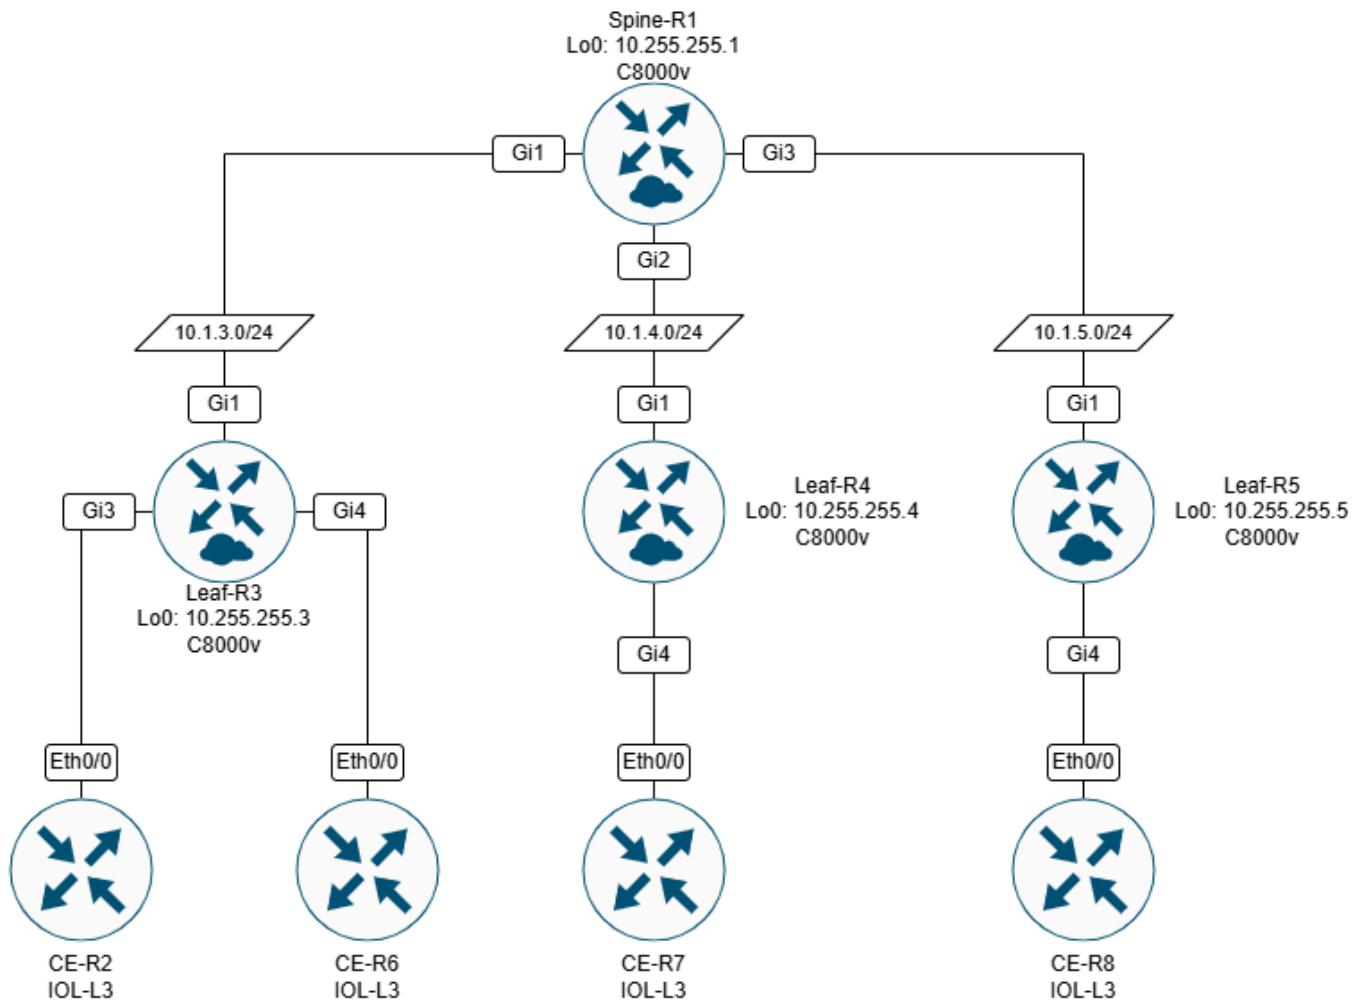


Cisco IOS-XE VxLAN L2VNI EVPN



Использованные образы:

- i86bi_LinuxL3-AdvEnterpriseK9-M2_157_3_May_2018.bin
- c8000v-17-13-01a

1. Summary:

В данной лабораторной работе необходимо организовать L2-связность между CE-маршрутизаторами. Для этого необходимо организовать VxLAN L2VNI при помощи EVPN.

1. Solution config:

Исходим из того, что стартовая конфигурация (адресация, IGP, PIM) уже применена.

В данной лабораторной R1 – это транзитный маршрутизатор, R3, R4 и R5 – это VTEP'ы.

В предыдущей лабораторной (VxLAN L2VNI Flood and Learn) MAC-адреса изучались на уровне data plane как при CE-VTEP-взаимодействии, так и при VTEP-VTEP-взаимодействии. Грубо говоря, о том, что на R3 появился какой-то MAC-адрес остальные VTEP'ы узнавали только после фактического получения какого-либо VxLAN-пакета с Ethernet-кадром, в котором, в свою очередь, в SRC_ADDR был указан этот MAC-адрес.

На этот раз мы будем использовать EVPN в качестве механизма передачи информации об изученных на VTEP'e MAC-адресах. В чем разница? Разница в том, что в этом случае VTEP посредством EVPN просигнализирует другим VTEP об изученном MAC-адресе, без необходимости отправки кадра к ним.

EVPN – Ethernet VPN – это относительно новая технология, устраниющая недостатки традиционных VPLS (невозможность полноценного мультихоминга, изучение адресов только на уровне data plane и т.п.). Как было сказано выше, EVPN позволяет просигнализировать другим VTEP изученные MAC-адреса (и не только). Сигнализация происходит при помощи BGP, AFI: L2VPN (25), SAFI: EVPN (70).

Итак, для настройки сервисов нам необходимо будет выполнить следующие шаги:

- настроить CE-маршрутизаторы
- настроить NVE-интерфейс
- настроить BGP-связность между VTEP'ами
- настроить L2-сервисы на VTEP'ах, а именно:
 - EFP (Ethernet Flow Point)
 - Bridge-domain
 - EVI (EVPN Instance)
 - VNI (VxLAN Network Identifier)

Начнем с настройки CE: для сервиса будем использовать влан 1001.

CE-R2	CE-R6
<pre>! interface Ethernet0/0.1001 encapsulation dot1Q 1001 ip address 10.10.1.2 255.255.255.0 !</pre>	<pre>! interface Ethernet0/0.1001 encapsulation dot1Q 1001 ip address 10.10.1.6 255.255.255.0 !</pre>
CE-R7	CE-R8
<pre>! interface Ethernet0/0.1001 encapsulation dot1Q 1001 ip address 10.10.1.7 255.255.255.0 !</pre>	<pre>! interface Ethernet0/0.1001 encapsulation dot1Q 1001 ip address 10.10.1.8 255.255.255.0 !</pre>

Теперь создадим NVE-интерфейс на каждом VTEP:

R3, R4, R5:

```
interface nve1
  source-interface Loopback0
!
```

Теперь давайте создадим 2 EFP для Leaf-R3 и добавим их в bridge-domain:

```
!
interface GigabitEthernet3
service instance 1001 ethernet
  encapsulation dot1q 1001
  rewrite ingress tag pop 1 symmetric
```

```

!
interface GigabitEthernet4
service instance 1001 ethernet
  encapsulation dot1q 1001
  rewrite ingress tag pop 1 symmetric
!
bridge-domain 1001
  member GigabitEthernet3 service-instance 1001
  member GigabitEthernet4 service-instance 1001
!
```

И проверим связность:

```

CE-R2#ping 10.10.1.6
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 10.10.1.6, timeout is 2 seconds:
.!!!!
Success rate is 80 percent (4/5), round-trip min/avg/max = 1/1/1 ms
CE-R2#
```

Для обеспечения юникастовой связности настроим на R2 и R6 BGP (никакой маршрутной информации там передаваться не будет – только KEEPALIVE-сообщения):

CE-R2	CE-R6
<pre> ! router bgp 65001 bgp log-neighbor-changes neighbor 10.10.1.6 remote-as 65001 !</pre>	<pre> ! router bgp 65001 bgp log-neighbor-changes neighbor 10.10.1.2 remote-as 65001 !</pre>

Оставим пока в покое CE-устройства и настроим BGP-связность в L2VPN EVPN.

В рамках данной лабораторной работы мы будем использовать R1 в качестве route-reflector, а Leaf'ы в свою очередь будут подключаться к нему как rr-клиенты.

Spine-R1	Leaf-R3 / R4 / R5
<pre> ! router bgp 65000 bgp log-neighbor-changes no bgp default ipv4-unicast neighbor RRC peer-group neighbor RRC remote-as 65000 neighbor RRC update-source Loopback0 neighbor 10.255.255.3 peer-group RRC neighbor 10.255.255.4 peer-group RRC neighbor 10.255.255.5 peer-group RRC ! address-family ipv4 exit-address-family ! address-family l2vpn evpn neighbor RRC send-community both neighbor RRC route-reflector-client neighbor 10.255.255.3 activate neighbor 10.255.255.4 activate neighbor 10.255.255.5 activate exit-address-family !</pre>	<pre> ! router bgp 65000 bgp log-neighbor-changes no bgp default ipv4-unicast neighbor 10.255.255.1 remote-as 65000 neighbor 10.255.255.1 update-source Loopback0 ! address-family ipv4 exit-address-family ! address-family l2vpn evpn neighbor 10.255.255.1 activate neighbor 10.255.255.1 send-community both exit-address-family !</pre>

Проверяем:

```

Spine-R1#show bgp l2vpn evpn summary
BGP router identifier 10.255.255.1, local AS number 65000
BGP table version is 1, main routing table version 1
```

Neighbor	V	AS	MsgRcvd	MsgSent	TblVer	InQ	OutQ	Up/Down	State/PfxRcd
10.255.255.3	4	65000	2	2	1	0	0	00:00:11	0
10.255.255.4	4	65000	5	5	1	0	0	00:00:57	0

```
10.255.255.5      4          65000        2          2          1          0          0 00:00:32          0
Spine-R1#
```

```
Leaf-R3#show bgp l2vpn evpn summary
BGP router identifier 10.255.255.3, local AS number 65000
BGP table version is 1, main routing table version 1
```

Neighbor	V	AS	MsgRcvd	MsgSent	TblVer	InQ	OutQ	Up/Down	State/PfxRcd
10.255.255.1	4	65000	2	2	1	0	0	00:00:39	0

```
Leaf-R3#show bgp l2vpn evpn
```

```
Leaf-R3#
```

BGP-сессии поднялись, в таблицах, как видно, пока пусто.

Теперь перейдем к настройкам EVPN на R3. Сначала глобальные настройки:

```
Leaf-R3(config)#l2vpn evpn
Leaf-R3(config-evpn)#?
L2VPN EVPN global configuration commands:
 default           Set a command to its defaults
 default-gateway   Default Gateway parameters
 exit              Exit from L2VPN evpn global configuration mode
 flooding-suppression Suppress flooding of broadcast, multicast, and/or
                        unknown unicast packets
 ip                IP parameters
 logging           Configure logging flags
 mac               MAC parameters
 mpls              MPLS parameters
 multihoming       Multihoming parameters
 no                Negate a command or set its defaults
 replication-type  Specify method for replicating BUM traffic
 route-target      Route Target VPN Extended Communities
 router-id         EVPN router ID
```

```
Leaf-R3(config-evpn)#

```

На данный момент нас здесь будут интересовать только следующие параметры:

1. **router-id** – здесь мы в явном виде укажем на основании адресации какого интерфейса будет создаваться **route-distinguisher** для наших EVI.
2. **replication-type** – здесь мы выберем метод отправки BUM-трафика. Нам доступны два варианта:
 - a. **ingress** – речь про ingress replication, т.е. отправку юникастом BUM-трафика в сторону каждого прописанного руками пира;
 - b. **static** – неочевидно, но здесь идет речь про отправку BUM-трафика через multicast-группы.

В явном виде укажем, что RD для всех EVI будет генерироваться на основании Loopback0, а BUM-трафик пойдет через мультикаст.

```
!
l2vpn evpn
 replication-type static
 router-id Loopback0
!
```

Теперь создадим EVI:

```
Leaf-R3(config)#l2vpn evpn instance 1001 vlan-based
Leaf-R3(config-evpn-evi)#encapsulation vxlan
Leaf-R3(config-evpn-evi)#ip local-learning disable
```

Что мы здесь сделали? Мы создали EVI под номером 1001 типа **vlan-based** (грубо говоря: один влан – один EVI – одна бридж-таблица, см. [RFC7432#6.1](#) для подробностей). Мы указали, что инкапсулироваться пакеты должны в VxLAN и в явном виде запретили изучение IP-адресов.

Посмотрим, что сейчас видно в EVI и включим debug исходящих update'ов.:

```

Leaf-R3#show l2vpn evpn evi 1001 detail
EVPN instance:          1001 (VLAN Based)
  RD:                  10.255.255.3:1001 (auto)
  Import-RTs:           65000:1001
  Export-RTs:           65000:1001
  Per-EVI Label:       none
  State:               Established
  Replication Type:   Static (global)
  Encapsulation:      vxlan
  IP Local Learn:     Disabled
  Adv. Def. Gateway:  Disabled (global)
  Re-originate RT5:    Disabled
  AR Flood Suppress:  Enabled (global)

```

```

Leaf-R3#show bgp l2vpn evpn
Leaf-R3#
Leaf-R3#debug bgp l2vpn evpn updates out
BGP updates debugging is on (outbound) for address family: L2VPN E-VPN

```

Теперь посмотрим, что есть в бриdge и добавим туда EVI:

```

Leaf-R3#show bridge-domain 1001
Bridge-domain 1001 (2 ports in all)
State: UP                   Mac learning: Enabled
Aging-Timer: 300 second(s)
Unknown Unicast Flooding Suppression: Disabled
Maximum address limit: 65536
  GigabitEthernet3 service instance 1001
  GigabitEthernet4 service instance 1001
  AED MAC address      Policy   Tag      Age  Pseudoport
  -----
  0  AABB.CC00.6000  forward dynamic  297  GigabitEthernet4.EFP1001
  0  AABB.CC00.2000  forward dynamic  298  GigabitEthernet3.EFP1001

```

```

Leaf-R3#
Leaf-R3#conf t
Enter configuration commands, one per line. End with CNTL/Z.
Leaf-R3(config)#bridge-domain 1001
Leaf-R3(config-bdomain)#member evpn-instance 1001 vni 10001
Leaf-R3(config-bdomain)#end
Leaf-R3#
*May 23 2025 19:41:51.435 MSK: %SYS-5-CONFIG_I: Configured from console by console
Leaf-R3#show bridge-domain 1001
*May 23 2025 19:42:03.793 MSK: BGP: EVPN Rcvd pfx: [2][10.255.255.3:1001][0][48][AABBCC006000][0][*]/20,
net flags: 0
*May 23 2025 19:42:03.793 MSK: TRM communities not added to sourced RT2
*May 23 2025 19:42:03.793 MSK: BGP(10): update modified for
[2][10.255.255.3:1001][0][48][AABBCC006000][0][*]/23
*May 23 2025 19:42:03.794 MSK: BGP(10): 10.255.255.1 Path-gateway-IPv4: IPv4 NEXT_HOP set to vxlan local
vtep-ip UNKNOWN for local net [2][10.255.255.3:1001][0][48][AABBCC006000][0][*]/20, nh_size: 4
*May 23 2025 19:42:03.794 MSK: BGP(10): update modified for
[2][10.255.255.3:1001][0][48][AABBCC006000][0][*]/23
*May 23 2025 19:42:03.794 MSK: BGP(10): (base) 10.255.255.1 send UPDATE (format)
[2][10.255.255.3:1001][0][48][AABBCC006000][0][*]/23, next 0.0.0.0, metric 0, path Local, extended
community RT:65000:1001 ENCAP:8
*May 23 2025 19:42:03.997 MSK: BGP: EVPN Rcvd pfx: [2][10.255.255.3:1001][0][48][AABBCC002000][0][*]/20,
net flags: 0
*May 23 2025 19:42:03.997 MSK: TRM communities not added to sourced RT2
*May 23 2025 19:42:03.997 MSK: BGP(10): update modified for
[2][10.255.255.3:1001][0][48][AABBCC002000][0][*]/23
Leaf-R3#conf t
*May 23 2025 19:42:03.997 MSK: BGP(10): 10.255.255.1 Path-gateway-IPv4: IPv4 NEXT_HOP set to vxlan local
vtep-ip UNKNOWN for local net [2][10.255.255.3:1001][0][48][AABBCC002000][0][*]/20, nh_size: 4
*May 23 2025 19:42:03.997 MSK: BGP(10): update modified for
[2][10.255.255.3:1001][0][48][AABBCC002000][0][*]/23
*May 23 2025 19:42:03.997 MSK: BGP(10): (base) 10.255.255.1 send UPDATE (format)
[2][10.255.255.3:1001][0][48][AABBCC002000][0][*]/23, next 0.0.0.0, metric 0, path Local, extended
community RT:65000:1001 ENCAP:8
Leaf-R3#

```

В дебаге мы видим, что в сторону R1 (RR) отправляется Update с двумя маршрутами типа 2, содержащих в себе RD (10.255.255.3:1001), RT (65000:1001), изученные MAC-адреса (по одному на каждый маршрут) и тип инкапсуляции (8 - VXLAN).

На скрине ниже содержимое этого пакета:

990 2025-05-22 22:04:44,653180	10.255.255.3	10.255.255.1	BGP	164 UPDATE Message
992 2025-05-22 22:04:44,854704	10.255.255.3	10.255.255.1	BGP	164 UPDATE Message
<hr/>				
> Transmission Control Protocol, Src Port: 179, Dst Port: 14689, Seq: 822, Ack: 855, Len: 110				
✓ Border Gateway Protocol - UPDATE Message				
Marker: ffffffffffffffffffffff				
Length: 110				
Type: UPDATE Message (2)				
Withdrawn Routes Length: 0				
Total Path Attribute Length: 87				
▼ Path attributes				
▼ Path Attribute - MP_REACH_NLRI				
> Flags: 0x80, Optional, Non-transitive, Complete				
Type Code: MP_REACH_NLRI (14)				
Length: 44				
Address family identifier (AFI): Layer-2 VPN (25)				
Subsequent address family identifier (SAFI): EVPN (70)				
> Next hop: 0.0.0.0				
Number of Subnetwork points of attachment (SNPA): 0				
▼ Network Layer Reachability Information (NLRI)				
▼ EVPN NLRI: MAC Advertisement Route				
Route Type: MAC Advertisement Route (2)				
Length: 33				
Route Distinguisher: 00010affff0303e9 (10.255.255.3:1001)				
> ESI: 00:00:00:00:00:00:00:00:00:00:00:00				
Ethernet Tag ID: 0				
MAC Address Length: 48				
MAC Address: aa:bb:cc:00:20:00				
IP Address Length: 0				
> IP Address: NOT INCLUDED				
VNI: 10001				
> Path Attribute - ORIGIN: INCOMPLETE				
> Path Attribute - AS_PATH: empty				
> Path Attribute - MULTI_EXIT_DISC: 0				
> Path Attribute - LOCAL_PREF: 100				
▼ Path Attribute - EXTENDED_COMMUNITIES				
> Flags: 0xc0, Optional, Transitive, Complete				
Type Code: EXTENDED_COMMUNITIES (16)				
Length: 16				
▼ Carried extended communities: (2 communities)				
▼ Route Target: 65000:1001 [Transitive 2-Octet AS-Specific]				
> Type: Transitive 2-Octet AS-Specific (0x00)				
Subtype (AS2): Route Target (0x02)				
2-Octet AS: 65000				
> Local Administrator: 0x00, Type: VID (802.1Q VLAN ID)				
Service Id: 1001				
▼ Encapsulation: VXLAN Encapsulation [Transitive Opaque]				
> Type: Transitive Opaque (0x03)				
Subtype (Opaque): Encapsulation (0x0c)				
Tunnel type: VXLAN Encapsulation (8)				

При этом на R1 мы увидим такое:

```
Spine-R1#
*May 22 2025 22:04:44.636 MSK: %BGP-6-NEXTHOP: Invalid next hop (0.0.0.0) received from 10.255.255.3:
martian next hop
Spine-R1#show bgp l2vpn evpn
Spine-R1#
```

Т.е. R1 из-за некорректного nexthop отбрасывает получаемые от R3 маршруты.

Чтобы разобраться в чем дело, вернемся на R3 и посмотрим внимательно на маршрутную информацию:

```
Leaf-R3#undbg all
All possible debugging has been turned off
Leaf-R3#
Leaf-R3#show bgp l2vpn evpn
BGP table version is 7, local router ID is 10.255.255.3
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
```

```

r RIB-failure, S Stale, m multipath, b backup-path, f RT-Filter,
x best-external, a additional-path, c RIB-compressed,
t secondary path, L long-lived-stale,
Origin codes: i - IGP, e - EGP, ? - incomplete
RPKI validation codes: V valid, I invalid, N Not found

```

Network	Next Hop	Metric	LocPrf	Weight	Path
Route Distinguisher: 10.255.255.3:1001					
> [2][10.255.255.3:1001][0][48][AABBCC002000][0][]/20	::	32768	?		
> [2][10.255.255.3:1001][0][48][AABBCC006000][0][]/20	::	32768	?		

Leaf-R3#show bgp l2vpn evpn detail

```

Route Distinguisher: 10.255.255.3:1001
BGP routing table entry for [2][10.255.255.3:1001][0][48][AABBCC002000][0][*]/20, version 7
  Paths: (1 available, best #1, table evi_1001)
    Advertised to update-groups:
      1
    Refresh Epoch 1
  Local
    :: (via default) from 0.0.0.0 (10.255.255.3)
      Origin incomplete, localpref 100, weight 32768, valid, sourced, local, best
      EVPN ESI: 000000000000000000000000, Label1 10001
      Extended Community: RT:65000:1001 ENCAP:8
      Local irb vxlan vtep:(inaccessible)
        rx pathid: 0, tx pathid: 0x0
        Updated on May 23 2025 19:42:03 MSK
BGP routing table entry for [2][10.255.255.3:1001][0][48][AABBCC006000][0][*]/20, version 6
  Paths: (1 available, best #1, table evi_1001)
    Advertised to update-groups:
      1
    Refresh Epoch 1
  Local
    :: (via default) from 0.0.0.0 (10.255.255.3)
      Origin incomplete, localpref 100, weight 32768, valid, sourced, local, best
      EVPN ESI: 000000000000000000000000, Label1 10001
      Extended Community: RT:65000:1001 ENCAP:8
      Local irb vxlan vtep:(inaccessible)
        rx pathid: 0, tx pathid: 0x0
        Updated on May 23 2025 19:42:03 MSK

```

Leaf-R3#

Желтым выделена проблема: у нас не настроен VNI. Собственно, именно это мы сейчас и сделаем:

```

Leaf-R3#debug bgp l2vpn evpn updates out
BGP updates debugging is on (outbound) for address family: L2VPN E-VPN
Leaf-R3#conf t
Enter configuration commands, one per line.  End with CNTL/Z.
Leaf-R3(config)#interface nve1
Leaf-R3(config-if)#host-reachability protocol bgp
Leaf-R3(config-if)#member vni 10001 mcast-group 239.1.0.1
Leaf-R3(config-if)#
*May 23 2025 19:56:05.540 MSK: TRM communities not added to sourced RT2
*May 23 2025 19:56:05.540 MSK: BGP(10): prefix [2][10.255.255.3:1001][0][48][AABBCC002000][0][*]/20
replacing esi starting 00 with 00 old vni label 002711 with new 002711 old vpn label 000000 with new
000000
*May 23 2025 19:56:05.540 MSK: TRM communities not added to sourced RT2
*May 23 2025 19:56:05.540 MSK: BGP(10): prefix [2][10.255.255.3:1001][0][48][AABBCC006000][0][*]/20
replacing esi starting 00 with 00 old vni label 002711 with new 002711 old vpn label 000000 with new
000000
*May 23 2025 19:56:05.541 MSK: BGP(10): update modified for
[2][10.255.255.3:1001][0][48][AABBCC002000][0][*]/23
*May 23 2025 19:56:05.541 MSK: BGP(10): 10.255.255.1 Path-gateway-IPv4: IPv4 NEXT_HOP set to vxlan local
vtep-ip 10.255.255.3 for local net [2][10.255.255.3:1001][0][48][AABBCC002000][0][*]/20, nh_size: 4
*May 23 2025 19:56:05.541 MSK: BGP(10): update modified for
[2][10.255.255.3:1001][0][48][AABBCC002000][0][*]/23
*May 23 2025 19:56:05.541 MSK: BGP(10): (base) 10.255.255.1 send UPDATE (format)
[2][10.255.255.3:1001][0][48][AABBCC002000][0][*]/23, next 10.255.255.3, metric 0, path Local, extended
community RT:65000:1001 ENCAP:8
*May 23 2025 19:56:05.541 MSK: BGP(10): update modified for
[2][10.255.255.3:1001][0][48][AABBCC006000][0][*]/23

```

```

*May 23 2025 19:56:05.541 MSK: BGP(10): 10.255.255.1 Path-gateway-IPv4: IPv4 NEXT_HOP set to vxlan local
vtep-ip 10.255.255.3 for local net [2][10.255.255.3:1001][0][48][AABBCC006000][0][*]/20, nh_size: 4
*May 23 2025 19:56:05.541 MSK: BGP(10): update modified for
[2][10.255.255.3:1001][0][48][AABBCC006000][0][*]/23
*May 23 2025 19:56:05.543 MSK: BGP(10): 10.255.255.1 rcv UPDATE w/ attr: nexthop 10.255.255.3, origin ?, localpref 100, metric 0, originator 10.255.255.3, clusterlist 10.255.255.1, merged path , AS_PATH , community , large community , extended community RT:65000:1001 ENCAP:8, SSA attribute
*May 23 2025 19:56:05.543 MSK: BGPSSA ssaccount is 0, Tunnel attribute
*May 23 2025 19:56:05.543 MSK: Tunnel encap type: 0, encap size: 0, Link-state attribute: {}, PrefixSid attribute:
*May 23 2025 19:56:05.543 MSK: BGP(10): 10.255.255.1 rcv UPDATE about
[2][10.255.255.3:1001][0][48][AABBCC002000][0][*]/23 -- DENIED due to: ORIGINATOR is us; MP_REACH NEXTHOP is our own address;
*May 23 2025 19:56:05.543 MSK: BGP(10): 10.255.255.1 rcv UPDATE about
[2][10.255.255.3:1001][0][48][AABBCC006000][0][*]/23 -- DENIED due to: ORIGINATOR is us; MP_REACH NEXTHOP is our own address;
*May 23 2025 19:56:06.977 MSK: %PIM-5-DRCHG: DR change from neighbor 0.0.0.0 to 10.255.255.3 on interface Tunnel0
Leaf-R3(config-if)#end
*May 23 2025 19:56:56.221 MSK: %SYS-5-CONFIG_I: Configured from console by console
Leaf-R3#undebug all
All possible debugging has been turned off
Leaf-R3#

```

Итак, что мы сделали и что получилось в итоге?

Сначала мы применили host-reachability protocol bgp в настройках интерфейса nve1 – данная команда «включает» EVPN для VxLAN-туннелей, работающих через этот интерфейс. И, разумеется, создали VNI 10001, с подпиской на группу 239.1.0.1.

В дебаге видно, что в качестве nexht-hop'а для сгенерированных маршрутов 2 типа был выставлен vtep-ip локального VTEP, т.е. 10.255.255.3. В таком виде они были отданы на RR.

Теперь посмотрим, что у нас видно в EVPN и VNI:

```

Leaf-R3#show l2vpn evpn evi 1001 detail
EVPN instance:          1001 (VLAN Based)
RD:                     10.255.255.3:1001 (auto)
Import-RTs:              65000:1001
Export-RTs:              65000:1001
Per-EVI Label:          none
State:                  Established
Replication Type:       Static (global)
Encapsulation:          vxlan
IP Local Learn:         Disabled
Adv. Def. Gateway:      Disabled (global)
Re-originate RT5:        Disabled
AR Flood Suppress:      Enabled (global)
Bridge Domain:          1001
Ethernet-Tag:           0
State:                  Established
Flood Suppress:          Detached
Core If:
Access If:
NVE If:                 nve1
RMAC:                   0000.0000.0000
Core BD:                 0
L2 VNI:                 10001
L3 VNI:                 0
VTEP IP:                10.255.255.3
MCAST IP:               239.1.0.1
Pseudoports:
  GigabitEthernet3 service instance 1001
    Routes: 1 MAC, 0 MAC/IP
  GigabitEthernet4 service instance 1001
    Routes: 1 MAC, 0 MAC/IP

```

```

Leaf-R3#show nve vni 10001
Interface  VNI      Multicast-group  VNI state  Mode   BD      cfg  vrf
nve1        10001    239.1.0.1       Up       L2CP   1001  CLI  N/A

```

Здесь мы видим что, evi 1001 привязан к BD 1001, к NVE-интерфейсу nve1, номер и тип VNI, адрес VTEP и мультикастовую группу. Также в выводе информации о VNI стоит обратить внимание на столбец Mode – значение в нем L2CP означает то, что это L2VNI и изучение адресной информации (читай MAC-адресов) происходит на CP (читай EVPN).

Теперь посмотрим, что у нас в BGP на R3 и R1:

```
Leaf-R3#show bgp 12vpn evpn AABBCC002000
BGP routing table entry for [2][10.255.255.3:1001][0][48][AABBCC002000][0][*]/20, version 18
Paths: (1 available, best #1, table evi_1001)
  Advertised to update-groups:
    1
  Refresh Epoch 1
  Local
    0.0.0.0 (via default) from 0.0.0.0 (10.255.255.3)
      Origin incomplete, localpref 100, weight 32768, valid, sourced, local, best
      EVPN ESI: 00000000000000000000, Label1 10001
      Extended Community: RT:65000:1001 ENCAP:8
        Local irb vxlan vtep:
          vrf:not found, l3-vni:0
          local router mac:0000.0000.0000
          core-irb interface:(not found)
          vtep-ip:10.255.255.3
        rx pathid: 0, tx pathid: 0x0
        Updated on May 23 2025 19:59:53 MSK
Leaf-R3#
```

```
Spine-R1#show bgp 12vpn evpn | b Network
  Network           Next Hop           Metric LocPrf Weight Path
Route Distinguisher: 10.255.255.3:1001
 *>i  [2][10.255.255.3:1001][0][48][AABBCC002000][0][*]/20
           10.255.255.3           0     100     0 ?
 *>i  [2][10.255.255.3:1001][0][48][AABBCC006000][0][*]/20
           10.255.255.3           0     100     0 ?
Spine-R1#show bgp 12vpn evpn AABBCC002000
BGP routing table entry for [2][10.255.255.3:1001][0][48][AABBCC002000][0][*]/20, version 10
Paths: (1 available, best #1, table EVPN-BGP-Table)
  Advertised to update-groups:
    1
  Refresh Epoch 1
  Local, (Received from a RR-client)
    10.255.255.3 (metric 10) (via default) from 10.255.255.3 (10.255.255.3)
      Origin incomplete, metric 0, localpref 100, valid, internal, best
      EVPN ESI: 00000000000000000000, Label1 10001
      Extended Community: RT:65000:1001 ENCAP:8
      rx pathid: 0, tx pathid: 0x0
      Updated on May 23 2025 19:59:53 MSK
Spine-R1#
```

Здесь мы видим, что информация о VTEP'е (выделено желтым) в общем-то является локальными данными, из которых, впрочем, как мы видели в дебаге, подставляются во время UPDATE'а значения (IP-адрес VTEP в качестве next-hop'a).

В поле Label1 вместо номера MPLS-метки передается (ожидаемо) номер VNI.

Теперь идем на R4 и R5, снова настраивать EVI (конфиг для обоих роутеров будет одинаковый):

```
12vpn evpn
 replication-type static
 router-id Loopback0
!
12vpn evpn instance 1001 vlan-based
 encapsulation vxlan
 ip local-learning disable
!
```

Сначала идем на R4. Но для начала посмотрим, видно ли что-нибудь сейчас в BGP:

```
Leaf-R4#show bgp 12vpn evpn
Leaf-R4#
```

```

Leaf-R4#conf t
Enter configuration commands, one per line. End with CNTL/Z.
Leaf-R4(config)#l2vpn evpn
Leaf-R4(config-evpn)# replication-type static
Leaf-R4(config-evpn)# router-id Loopback0
Leaf-R4(config-evpn)#
Leaf-R4(config-evpn)#l2vpn evpn instance 1001 vlan-based
Leaf-R4(config-evpn-evi)# encapsulation vxlan
Leaf-R4(config-evpn-evi)# ip local-learning disable
Leaf-R4(config-evpn-evi)#
Leaf-R4(config-evpn-evi)#end
Leaf-R4#

```

И посмотрим, изменилось ли что-нибудь:

```

Leaf-R4#show bgp l2vpn evpn | b Network
      Network          Next Hop          Metric LocPrf Weight Path
Route Distinguisher: 10.255.255.3:1001
 *>i [2][10.255.255.3:1001][0][48][AABBCC002000][0][*]/20
           10.255.255.3          0    100      0 ?
 *>i [2][10.255.255.3:1001][0][48][AABBCC006000][0][*]/20
           10.255.255.3          0    100      0 ?
Route Distinguisher: 10.255.255.4:1001
 *>i [2][10.255.255.4:1001][0][48][AABBCC002000][0][*]/20
           10.255.255.3          0    100      0 ?
 *>i [2][10.255.255.4:1001][0][48][AABBCC006000][0][*]/20
           10.255.255.3          0    100      0 ?

```

Leaf-R4#

No.	Time	Source	Destination	Protocol	Lengtl	Info
13	2025-05-23 20:37:22,610914	10.255.255.4	10.255.255.1	BGP	77	ROUTE-REFRESH Message
14	2025-05-23 20:37:22,611644	10.255.255.1	10.255.255.4	BGP	77	ROUTE-REFRESH Message
15	2025-05-23 20:37:22,611862	10.255.255.1	10.255.255.4	BGP	236	UPDATE Message, ROUTE-REFRESH Message
37	2025-05-23 20:38:16,686515	10.255.255.1	10.255.255.4	BGP	73	KEEPALIVE Message
40	2025-05-23 20:38:19,094047	10.255.255.4	10.255.255.1	BGP	73	KEEPALIVE Message

```

> Frame 15: 236 bytes on wire (1888 bits), 236 bytes captured (1888 bits) on interface -, id 0
> Ethernet II, Src: 50:00:00:01:00:01, Dst: 50:00:00:04:00:00
> Internet Protocol Version 4, Src: 10.255.255.1, Dst: 10.255.255.4
> Transmission Control Protocol, Src Port: 38227, Dst Port: 179, Seq: 43, Ack: 43, Len: 182
  Border Gateway Protocol - UPDATE Message
    Marker: ffffffffffffffffffffff
    Length: 159
    Type: UPDATE Message (2)
    Withdrawn Routes Length: 0
    Total Path Attribute Length: 136
  Path attributes
    Path Attribute - MP_REACH_NLRI
      Flags: 0x80, Optional, Non-transitive, Complete
      Type Code: MP_REACH_NLRI (14)
      Length: 79
      Address family identifier (AFI): Layer-2 VPN (25)
      Subsequent address family identifier (SAFI): EVPN (70)
      Next hop: 10.255.255.3
      Number of Subnetwork points of attachment (SNPA): 0
    Network Layer Reachability Information (NLRI)
      EVPN NLRI: MAC Advertisement Route
        Route Type: MAC Advertisement Route (2)
        Length: 33
        Route Distinguisher: 00010affff0303e9 (10.255.255.3:1001)
        ESI: 00:00:00:00:00:00:00:00:00:00:00:00:00:00:00:00:00:00
        Ethernet Tag ID: 0
        MAC Address Length: 48
        MAC Address: aa:bb:cc:00:20:00
        IP Address Length: 0
        IP Address: NOT INCLUDED
        VNI: 10001
      EVPN NLRI: MAC Advertisement Route
    Path Attribute - ORIGIN: INCOMPLETE
    Path Attribute - AS_PATH: empty
    Path Attribute - MULTI_EXIT_DISC: 0
    Path Attribute - LOCAL_PREF: 100
  Path Attribute - EXTENDED_COMMUNITIES
    Flags: 0xc0, Optional, Transitive, Complete
    Type Code: EXTENDED_COMMUNITIES (16)
    Length: 16
    Carried extended communities: (2 communities)
      Route Target: 65000:1001 [Transitive 2-Octet AS-Specific]
      Encapsulation: VXLAN Encapsulation [Transitive Opaque]
    Path Attribute - CLUSTER_LIST: 10.255.255.1
    Path Attribute - ORIGINATOR_ID: 10.255.255.3
  Border Gateway Protocol - ROUTE-REFRESH Message

```

Что мы здесь видим? А видим мы здесь то, что как только мы создали EVI 1001, R4 отправил запрос ROUTE-REFRESH в сторону рефлектора и, в итоге, получил UPDATE с маршрутной информацией (см. скрин).

Важно отметить, что сейчас мы даже не успели создать BD и, тем более, привязать к нему EVI.

Сейчас в таблице BGP EVPN находится 4 маршрута: 2 с RD 10.255.255.3:1001 (полученные от RR) и 2 этих же маршрута, но уже «импортированные» в таблицу evi_1001 с соответствующим RD.

Если посмотреть на EVI, то можно будет увидеть почему эти маршруты были импортированы в этот MAC-VRF: у нас совпадают значения Import-RT на R4 со значением Export-RT на R3.

Leaf-R4	Leaf-R3
<pre>Leaf-R4#show l2vpn evpn evi 1001 detail EVPN instance: 1001 (VLAN Based) RD: 10.255.255.4:1001 (auto) Import-RTs: 65000:1001 Export-RTs: 65000:1001 Per-EVI Label: none State: Established Replication Type: Static (global) Encapsulation: vxlan IP Local Learn: Disabled Adv. Def. Gateway: Disabled (global) Re-originate RT5: Disabled AR Flood Suppress: Enabled (global)</pre>	<pre>Leaf-R3#show l2vpn evpn evi 1001 detail EVPN instance: 1001 (VLAN Based) RD: 10.255.255.3:1001 (auto) Import-RTs: 65000:1001 Export-RTs: 65000:1001 Per-EVI Label: none State: Established Replication Type: Static (global) Encapsulation: vxlan IP Local Learn: Disabled Adv. Def. Gateway: Disabled (global) Re-originate RT5: Disabled AR Flood Suppress: Enabled (global) Bridge Domain: 1001 Ethernet-Tag: 0 State: Established Flood Suppress: Detached Core If: Access If: NVE If: nve1 RMAC: 0000.0000.0000 Core BD: 0 L2 VNI: 10001 L3 VNI: 0 VTEP IP: 10.255.255.3 MCAST IP: 239.1.0.1 Pseudoports: GigabitEthernet3 service instance 1001 Routes: 1 MAC, 0 MAC/IP GigabitEthernet4 service instance 1001 Routes: 1 MAC, 0 MAC/IP</pre>

Теперь посмотрим повнимательнее на сами маршруты, на примере MAC-адреса CE-R2:

```
Leaf-R4#show bgp l2vpn evpn AABBCC002000
BGP routing table entry for [2][10.255.255.3:1001][0][48][AABBCC002000][0][*]/20, version 18
Paths: (1 available, best #1, table EVPN-BGP-Table)
  Not advertised to any peer
  Refresh Epoch 6
  Local
    10.255.255.3 (metric 20) (via default) from 10.255.255.1 (10.255.255.1)
      Origin incomplete, metric 0, localpref 100, valid, internal, best
      EVPN ESI: 0000000000000000, Label1 10001
      Extended Community: RT:65000:1001 ENCAP:8
      Originator: 10.255.255.3, Cluster list: 10.255.255.1
      rx pathid: 0, tx pathid: 0x0
      Updated on May 23 2025 20:37:22 MSK
BGP routing table entry for [2][10.255.255.4:1001][0][48][AABBCC002000][0][*]/20, version 20
Paths: (1 available, best #1, table evi_1001)
  Not advertised to any peer
  Refresh Epoch 6
  Local, imported path from [2][10.255.255.3:1001][0][48][AABBCC002000][0][*]/20 (global)
    10.255.255.3 (metric 20) (via default) from 10.255.255.1 (10.255.255.1)
      Origin incomplete, metric 0, localpref 100, valid, internal, best
      EVPN ESI: 0000000000000000, Label1 10001
      Extended Community: RT:65000:1001 ENCAP:8
      Originator: 10.255.255.3, Cluster list: 10.255.255.1
      Local irb vxlan vtep:(inaccessible)
      rx pathid: 0, tx pathid: 0x0
      Updated on May 23 2025 20:37:22 MSK
```

Здесь мы в явном виде увидим номер VNI, откуда и куда импортированы маршруты, номер VTEP'а и т.п.

Создадим bridge-domain, EFP и донастроим интерфейс NVE:

```
!
interface Gi4
  service instance 1001 ethernet
    encapsulation dot1q 1001
  !
!
interface nve1
  host-reachability protocol bgp
  member vni 10001 mcast-group 239.1.0.1
!
bridge-domain 1001
  member GigabitEthernet4 service-instance 1001
  member evpn-instance 1001 vni 10001
!
```

И проверим:

```
Leaf-R4#show bridge-domain 1001
Bridge-domain 1001 (2 ports in all)
State: UP          Mac learning: Enabled
Aging-Timer: 300 second(s)
Unknown Unicast Flooding Suppression: Disabled
Maximum address limit: 65536
  GigabitEthernet4 service instance 1001
  vni 10001
  AED MAC address   Policy   Tag      Age  Pseudoport
  -----
  -  AABB.CC00.6000 forward static_r 0    nve1.VNI10001, EVPN
  -  AABB.CC00.2000 forward static_r 0    nve1.VNI10001, EVPN

Leaf-R4#show l2vpn evpn mac evi 1001
MAC Address   EVI   BD   ESI           Ether Tag  Next Hop(s)
  -----
aabb.cc00.2000 1001  1001  0000.0000.0000.0000.0000 0        10.255.255.3
aabb.cc00.6000 1001  1001  0000.0000.0000.0000.0000 0        10.255.255.3

Leaf-R4#show l2vpn evpn mac evi 1001 detail
MAC Address:          aabb.cc00.2000
EVPN Instance:        1001
Bridge Domain:        1001
Ethernet Segment:    0000.0000.0000.0000.0000
Ethernet Tag ID:     0
Next Hop(s):          V:10001 10.255.255.3
Local Address:        10.255.255.4
Sequence Number:      0
MAC only present:    Yes
MAC Duplication Detection: Timer not running

MAC Address:          aabb.cc00.6000
EVPN Instance:        1001
Bridge Domain:        1001
Ethernet Segment:    0000.0000.0000.0000.0000
Ethernet Tag ID:     0
Next Hop(s):          V:10001 10.255.255.3
Local Address:        10.255.255.4
Sequence Number:      0
MAC only present:    Yes
MAC Duplication Detection: Timer not running
```

В выводе BD мы видим, что таблице MAC-адресов есть две записи, полученные по EVPN и доступные через VNI 10001. Если мы повнимательнее посмотрим на MAC-адреса в l2vpn evpn, то сможем увидеть, какие маршрутизаторы являются next-hop'ами.

Попробуем пропинговать R2 и R6 с R7:

```

CE-R7#ping 10.10.1.2 rep 2
Type escape sequence to abort.
Sending 2, 100-byte ICMP Echos to 10.10.1.2, timeout is 2 seconds:
!
Success rate is 50 percent (1/2), round-trip min/avg/max = 2/2/2 ms
CE-R7#ping 10.10.1.6 rep 2
Type escape sequence to abort.
Sending 2, 100-byte ICMP Echos to 10.10.1.6, timeout is 2 seconds:
!
Success rate is 50 percent (1/2), round-trip min/avg/max = 1/1/1 ms
CE-R7#

```

Связность есть, а на скрине ниже видны все пакеты.

vxlan						
No.	Time	Source	Destination	Protocol	Length	Info
1348	2025-05-23 21:28:05,839385	aa:bb:cc:00:70:00	ff:ff:ff:ff:ff:ff	PIMv2	142	Register
1349	2025-05-23 21:28:05,840412	aa:bb:cc:00:70:00	ff:ff:ff:ff:ff:ff	ARP	114	Who has 10.10.1.2? Tell 10.10.1.7
1351	2025-05-23 21:28:05,843007	aa:bb:cc:00:20:00	aa:bb:cc:00:70:00	ARP	114	10.10.1.2 is at aa:bb:cc:00:20:00
1354	2025-05-23 21:28:07,797132	10.10.1.7	10.10.1.2	ICMP	168	Echo (ping) request id=0x0000, seq=1/256, ttl=255 (reply in 1355)
1355	2025-05-23 21:28:07,798305	10.10.1.2	10.10.1.7	ICMP	168	Echo (ping) reply id=0x0000, seq=1/256, ttl=255 (request in 1354)
1363	2025-05-23 21:28:25,427155	aa:bb:cc:00:70:00	ff:ff:ff:ff:ff:ff	PIMv2	142	Register
1364	2025-05-23 21:28:25,427167	aa:bb:cc:00:70:00	ff:ff:ff:ff:ff:ff	ARP	114	Who has 10.10.1.6? Tell 10.10.1.7
1365	2025-05-23 21:28:25,429791	aa:bb:cc:00:60:00	aa:bb:cc:00:70:00	ARP	114	10.10.1.6 is at aa:bb:cc:00:60:00
1368	2025-05-23 21:28:25,462326	aa:bb:cc:00:60:00	aa:bb:cc:00:70:00	ARP	114	10.10.1.6 is at aa:bb:cc:00:60:00
1369	2025-05-23 21:28:27,430024	10.10.1.7	10.10.1.6	ICMP	168	Echo (ping) request id=0x0001, seq=1/256, ttl=255 (reply in 1370)
1370	2025-05-23 21:28:27,430629	10.10.1.6	10.10.1.7	ICMP	168	Echo (ping) reply id=0x0001, seq=1/256, ttl=255 (request in 1369)

Wireshark - Пакет 1349 - Standard input

```
> Frame 1349: 114 bytes on wire (912 bits), 114 bytes captured (912 bits) on interface -, id 0
> Ethernet II, Src: 50:00:00:01:00:01, Dst: 01:00:00:01:00:01
> Internet Protocol Version 4, Src: 10.255.255.4, Dst: 239.1.0.1
> User Datagram Protocol, Src Port: 59459, Dst Port: 4789
> Virtual eXtensible Local Area Network
> Ethernet II, Src: aa:bb:cc:00:70:00, Dst: ff:ff:ff:ff:ff:ff
> 802.1Q Virtual LAN, PRI: 0, DEI: 0, ID: 1001
> Address Resolution Protocol (request)
```

Wireshark - Пакет 1351 - Standard input

```
> Frame 1351: 114 bytes on wire (912 bits), 114 bytes captured (912 bits) on interface -, id 1
> Ethernet II, Src: 50:00:00:01:00:01, Dst: 50:00:00:00:00:00
> Internet Protocol Version 4, Src: 10.255.255.3, Dst: 10.255.255.4
> User Datagram Protocol, Src Port: 58215, Dst Port: 4789
> Virtual eXtensible Local Area Network
> Ethernet II, Src: aa:bb:cc:00:70:00, Dst: aa:bb:cc:00:60:00
> 802.1Q Virtual LAN, PRI: 0, DEI: 0, ID: 1001
> Address Resolution Protocol (reply)
```

Теперь еще раз посмотрим, что у нас таблицах BGP:

```

Leaf-R4#show bgp 12vpn evpn | b Network
      Network          Next Hop          Metric LocPrf Weight Path
Route Distinguisher: 10.255.255.3:1001
  *>i [2][10.255.255.3:1001][0][48][AABBCC002000][0][*]/20
    10.255.255.3          0        100      0 ?
  *>i [2][10.255.255.3:1001][0][48][AABBCC006000][0][*]/20
    10.255.255.3          0        100      0 ?
Route Distinguisher: 10.255.255.4:1001
  *>i [2][10.255.255.4:1001][0][48][AABBCC002000][0][*]/20
    10.255.255.3          0        100      0 ?
  *>i [2][10.255.255.4:1001][0][48][AABBCC006000][0][*]/20
    10.255.255.3          0        100      0 ?
  *>  [2][10.255.255.4:1001][0][48][AABBCC007000][0][*]/20
    0.0.0.0                  32768 ?

```

Leaf-R4#

```

Leaf-R3# show bgp 12vpn evpn | b Network
      Network          Next Hop          Metric LocPrf Weight Path
Route Distinguisher: 10.255.255.3:1001
  *>  [2][10.255.255.3:1001][0][48][AABBCC002000][0][*]/20
    0.0.0.0                  32768 ?
  *>  [2][10.255.255.3:1001][0][48][AABBCC006000][0][*]/20
    0.0.0.0                  32768 ?
  *>i [2][10.255.255.3:1001][0][48][AABBCC007000][0][*]/20
    10.255.255.4          0        100      0 ?
Route Distinguisher: 10.255.255.4:1001
  *>i [2][10.255.255.4:1001][0][48][AABBCC007000][0][*]/20
    10.255.255.4          0        100      0 ?

```

Leaf-R3#

Теперь пришла пора настроить последний VTEP.

```
12vpn evpn
 replication-type static
```

```

router-id Loopback0
!
12vpn evpn instance 1001 vlan-based
  encapsulation vxlan
  ip local-learning disable
!
!
interface Gi4
  service instance 1001 ethernet
    encapsulation dot1q 1001
  !
!
interface nve1
  host-reachability protocol bgp
  member vni 10001 mcast-group 239.1.0.1
!
bridge-domain 1001
  member GigabitEthernet4 service-instance 1001
  member evpn-instance 1001 vni 10001
!

```

Применяем конфигурацию и проверяем:

```

Leaf-R5#show bgp 12vpn evpn | b Network
      Network          Next Hop          Metric LocPrf Weight Path
Route Distinguisher: 10.255.255.3:1001
  *>i [2][10.255.255.3:1001][0][48][AABBCC002000][0][*]/20
    10.255.255.3          0     100      0 ?
  *>i [2][10.255.255.3:1001][0][48][AABBCC006000][0][*]/20
    10.255.255.3          0     100      0 ?
Route Distinguisher: 10.255.255.4:1001
  *>i [2][10.255.255.4:1001][0][48][AABBCC007000][0][*]/20
    10.255.255.4          0     100      0 ?
Route Distinguisher: 10.255.255.5:1001
  *>i [2][10.255.255.5:1001][0][48][AABBCC002000][0][*]/20
    10.255.255.3          0     100      0 ?
  *>i [2][10.255.255.5:1001][0][48][AABBCC006000][0][*]/20
    10.255.255.3          0     100      0 ?
  *>i [2][10.255.255.5:1001][0][48][AABBCC007000][0][*]/20
    10.255.255.4          0     100      0 ?

Leaf-R5#
Leaf-R5#show bridge-domain 1001
Bridge-domain 1001 (2 ports in all)
State: UP          Mac learning: Enabled
Aging-Timer: 300 second(s)
Unknown Unicast Flooding Suppression: Disabled
Maximum address limit: 65536
  GigabitEthernet4 service instance 1001
    vni 10001
    AED MAC address   Policy   Tag      Age  Pseudopoport
    -----
    -  AABB.CC00.6000  forward static_r 0    nve1.VNI10001, EVPN
    -  AABB.CC00.7000  forward static_r 0    nve1.VNI10001, EVPN
    -  AABB.CC00.2000  forward static_r 0    nve1.VNI10001, EVPN

Leaf-R5#show 12vpn evpn mac evi 1001 detail
MAC Address:          aabb.cc00.2000
EVPN Instance:         1001
Bridge Domain:         1001
Ethernet Segment:     0000.0000.0000.0000.0000
Ethernet Tag ID:      0
Next Hop(s):          V:10001 10.255.255.3
Local Address:         10.255.255.5
Sequence Number:       0
MAC only present:     Yes
MAC Duplication Detection: Timer not running

MAC Address:          aabb.cc00.6000
EVPN Instance:         1001
Bridge Domain:         1001
Ethernet Segment:     0000.0000.0000.0000.0000
Ethernet Tag ID:      0

```

Next Hop(s): V:10001 10.255.255.3
 Local Address: 10.255.255.5
 Sequence Number: 0
 MAC only present: Yes
 MAC Duplication Detection: Timer not running

MAC Address: aabb.cc00.7000
 EVPN Instance: 1001
 Bridge Domain: 1001
 Ethernet Segment: 0000.0000.0000.0000.0000
 Ethernet Tag ID: 0
 Next Hop(s): V:10001 10.255.255.4
 Local Address: 10.255.255.5
 Sequence Number: 0
 MAC only present: Yes
 MAC Duplication Detection: Timer not running

Теперь давайте настроим на всех 4 CE-роутерах Full-mesh iBGP и, таким образом, проверим связность «каждый с каждым».

CE-R2	CE-R6
<pre>router bgp 65001 neighbor 10.10.1.6 remote-as 65001 neighbor 10.10.1.7 remote-as 65001 neighbor 10.10.1.8 remote-as 65001 !</pre>	<pre>router bgp 65001 neighbor 10.10.1.2 remote-as 65001 neighbor 10.10.1.7 remote-as 65001 neighbor 10.10.1.8 remote-as 65001 !</pre>
CE-R7	CE-R8
<pre>router bgp 65001 neighbor 10.10.1.2 remote-as 65001 neighbor 10.10.1.6 remote-as 65001 neighbor 10.10.1.8 remote-as 65001 !</pre>	<pre>router bgp 65001 neighbor 10.10.1.2 remote-as 65001 neighbor 10.10.1.6 remote-as 65001 neighbor 10.10.1.7 remote-as 65001 !</pre>

CE-R2#show bgp ipv4 unicast summary
 BGP router identifier 10.10.1.2, local AS number 65001
 BGP table version is 1, main routing table version 1

Neighbor	V	AS	MsgRcvd	MsgSent	TblVer	InQ	OutQ	Up/Down	State/PfxRcd
10.10.1.6	4	65001	4	6	1	0	0	00:01:48	0
10.10.1.7	4	65001	5	3	1	0	0	00:01:47	0
10.10.1.8	4	65001	4	4	1	0	0	00:01:47	0

CE-R6#show bgp ipv4 unicast summary
 BGP router identifier 10.10.1.6, local AS number 65001
 BGP table version is 1, main routing table version 1

Neighbor	V	AS	MsgRcvd	MsgSent	TblVer	InQ	OutQ	Up/Down	State/PfxRcd
10.10.1.2	4	65001	6	4	1	0	0	00:01:48	0
10.10.1.7	4	65001	3	3	1	0	0	00:01:45	0
10.10.1.8	4	65001	6	6	1	0	0	00:01:49	0

CE-R7#show bgp ipv4 unicast summary
 BGP router identifier 10.10.1.7, local AS number 65001
 BGP table version is 1, main routing table version 1

Neighbor	V	AS	MsgRcvd	MsgSent	TblVer	InQ	OutQ	Up/Down	State/PfxRcd
10.10.1.2	4	65001	3	5	1	0	0	00:01:47	0
10.10.1.6	4	65001	3	3	1	0	0	00:01:45	0
10.10.1.8	4	65001	3	3	1	0	0	00:01:42	0

CE-R8#show bgp ipv4 unicast summary
 BGP router identifier 10.10.1.8, local AS number 65001
 BGP table version is 1, main routing table version 1

Neighbor	V	AS	MsgRcvd	MsgSent	TblVer	InQ	OutQ	Up/Down	State/PfxRcd
10.10.1.2	4	65001	4	4	1	0	0	00:01:47	0
10.10.1.6	4	65001	6	6	1	0	0	00:01:49	0
10.10.1.7	4	65001	3	3	1	0	0	00:01:42	0

3. Solution config:

Spine-R1	Leaf-R3
<pre> router bgp 65000 bgp log-neighbor-changes no bgp default ipv4-unicast neighbor RRC peer-group neighbor RRC remote-as 65000 neighbor RRC update-source Loopback0 neighbor 10.255.255.3 peer-group RRC neighbor 10.255.255.4 peer-group RRC neighbor 10.255.255.5 peer-group RRC address-family l2vpn evpn neighbor RRC send-community both neighbor RRC route-reflector-client neighbor 10.255.255.3 activate neighbor 10.255.255.4 activate neighbor 10.255.255.5 activate ! </pre>	<pre> l2vpn evpn replication-type static router-id Loopback0 ! l2vpn evpn instance 1001 vlan-based encapsulation vxlan ip local-learning disable ! bridge-domain 1001 member GigabitEthernet3 service-instance 1001 member GigabitEthernet4 service-instance 1001 member evpn-instance 1001 vni 10001 ! interface GigabitEthernet3 service instance 1001 ethernet encapsulation dot1q 1001 ! interface GigabitEthernet4 service instance 1001 ethernet encapsulation dot1q 1001 ! interface nve1 source-interface Loopback0 host-reachability protocol bgp member vni 10001 mcast-group 239.1.0.1 ! router bgp 65000 bgp log-neighbor-changes no bgp default ipv4-unicast neighbor 10.255.255.1 remote-as 65000 neighbor 10.255.255.1 update-source Loopback0 address-family l2vpn evpn neighbor 10.255.255.1 activate neighbor 10.255.255.1 send-community both !</pre>
Leaf-R4	Leaf-R5
<pre> l2vpn evpn replication-type static router-id Loopback0 ! l2vpn evpn instance 1001 vlan-based encapsulation vxlan ip local-learning disable ! bridge-domain 1001 member GigabitEthernet4 service-instance 1001 member evpn-instance 1001 vni 10001 ! interface GigabitEthernet4 service instance 1001 ethernet encapsulation dot1q 1001 ! interface nve1 source-interface Loopback0 host-reachability protocol bgp member vni 10001 mcast-group 239.1.0.1 ! router bgp 65000 bgp log-neighbor-changes no bgp default ipv4-unicast neighbor 10.255.255.1 remote-as 65000 neighbor 10.255.255.1 update-source Loopback0 address-family l2vpn evpn neighbor 10.255.255.1 activate neighbor 10.255.255.1 send-community both !</pre>	<pre> l2vpn evpn replication-type static router-id Loopback0 ! l2vpn evpn instance 1001 vlan-based encapsulation vxlan ip local-learning disable ! bridge-domain 1001 member GigabitEthernet4 service-instance 1001 member evpn-instance 1001 vni 10001 ! interface GigabitEthernet4 service instance 1001 ethernet encapsulation dot1q 1001 ! interface nve1 source-interface Loopback0 host-reachability protocol bgp member vni 10001 mcast-group 239.1.0.1 ! router bgp 65000 bgp log-neighbor-changes no bgp default ipv4-unicast neighbor 10.255.255.1 remote-as 65000 neighbor 10.255.255.1 update-source Loopback0 address-family l2vpn evpn neighbor 10.255.255.1 activate neighbor 10.255.255.1 send-community both !</pre>

CE-R2	CE-R6
<pre>! interface Ethernet0/0.1001 encapsulation dot1Q 1001 ip address 10.10.1.2 255.255.255.0 ! router bgp 65001 neighbor 10.10.1.6 remote-as 65001 neighbor 10.10.1.7 remote-as 65001 neighbor 10.10.1.8 remote-as 65001 !</pre>	<pre>! interface Ethernet0/0.1001 encapsulation dot1Q 1001 ip address 10.10.1.6 255.255.255.0 ! router bgp 65001 neighbor 10.10.1.2 remote-as 65001 neighbor 10.10.1.7 remote-as 65001 neighbor 10.10.1.8 remote-as 65001 !</pre>
CE-R7	CE-R8
<pre>! interface Ethernet0/0.1001 encapsulation dot1Q 1001 ip address 10.10.1.7 255.255.255.0 ! router bgp 65001 neighbor 10.10.1.2 remote-as 65001 neighbor 10.10.1.6 remote-as 65001 neighbor 10.10.1.8 remote-as 65001 !</pre>	<pre>! interface Ethernet0/0.1001 encapsulation dot1Q 1001 ip address 10.10.1.8 255.255.255.0 ! router bgp 65001 neighbor 10.10.1.2 remote-as 65001 neighbor 10.10.1.6 remote-as 65001 neighbor 10.10.1.7 remote-as 65001 !</pre>